

# Flink 资源动态调整及其实实践

Self Regulating Stream Processing in Flink

马庆祥  
Qingxiang Ma

数据开发资深工程师  
Senior data development engineer

**FLINK FORWARD # ASIA**

实时即未来 # Real-time Is The Future

**FLINK  
FORWARD**



# Contents

## 目录

### 01 背景

Background

### 02 QDS 模型

The QDS model

### 03 生产实践

Production Practices

### 04 展望

Future Work

# 关于我们

About us



# 背景

Background

---

01



# 背景

Background



Stateful Computations over Data Streams



State

Savepoint

bug fixing

upgrades

reconfiguration

dynamic scaling

But, when and how much to scale?

# 背景

## Background

### 问题：

Problems:

◆ 大白同学：第一次用 Flink 进行业务开发，终于要上线了，如何配置并行度啊？

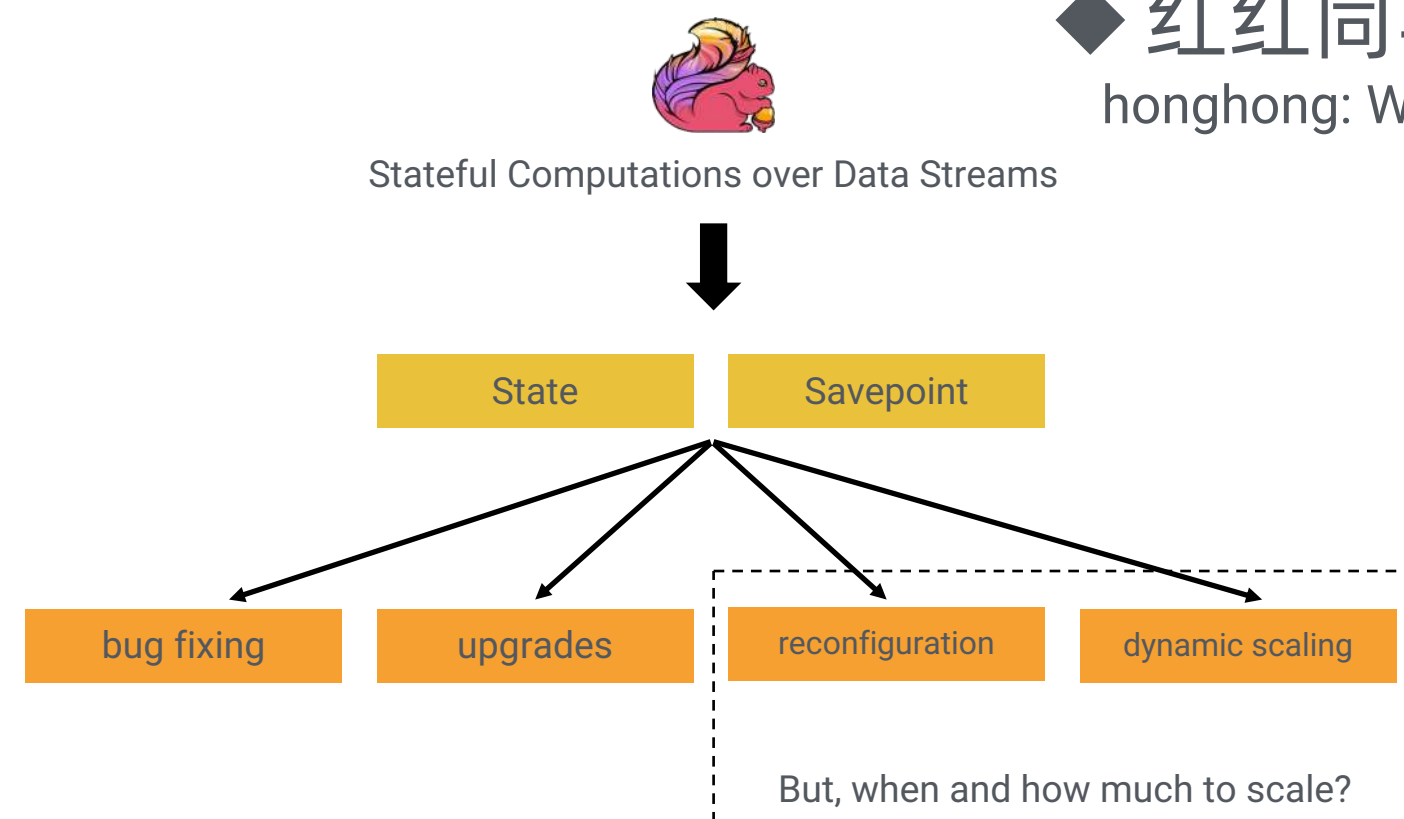
dabai: I am a newer of Flink, the job is coming online, how to configure parallelism?

◆ 小明同学：江湖救急，业务线数据流量突增，现在积压了大量的数据，该咋办呀？

xiaoming: Help, business line data traffic suddenly increased, and a large amount of data has been accumulated, what should I do?

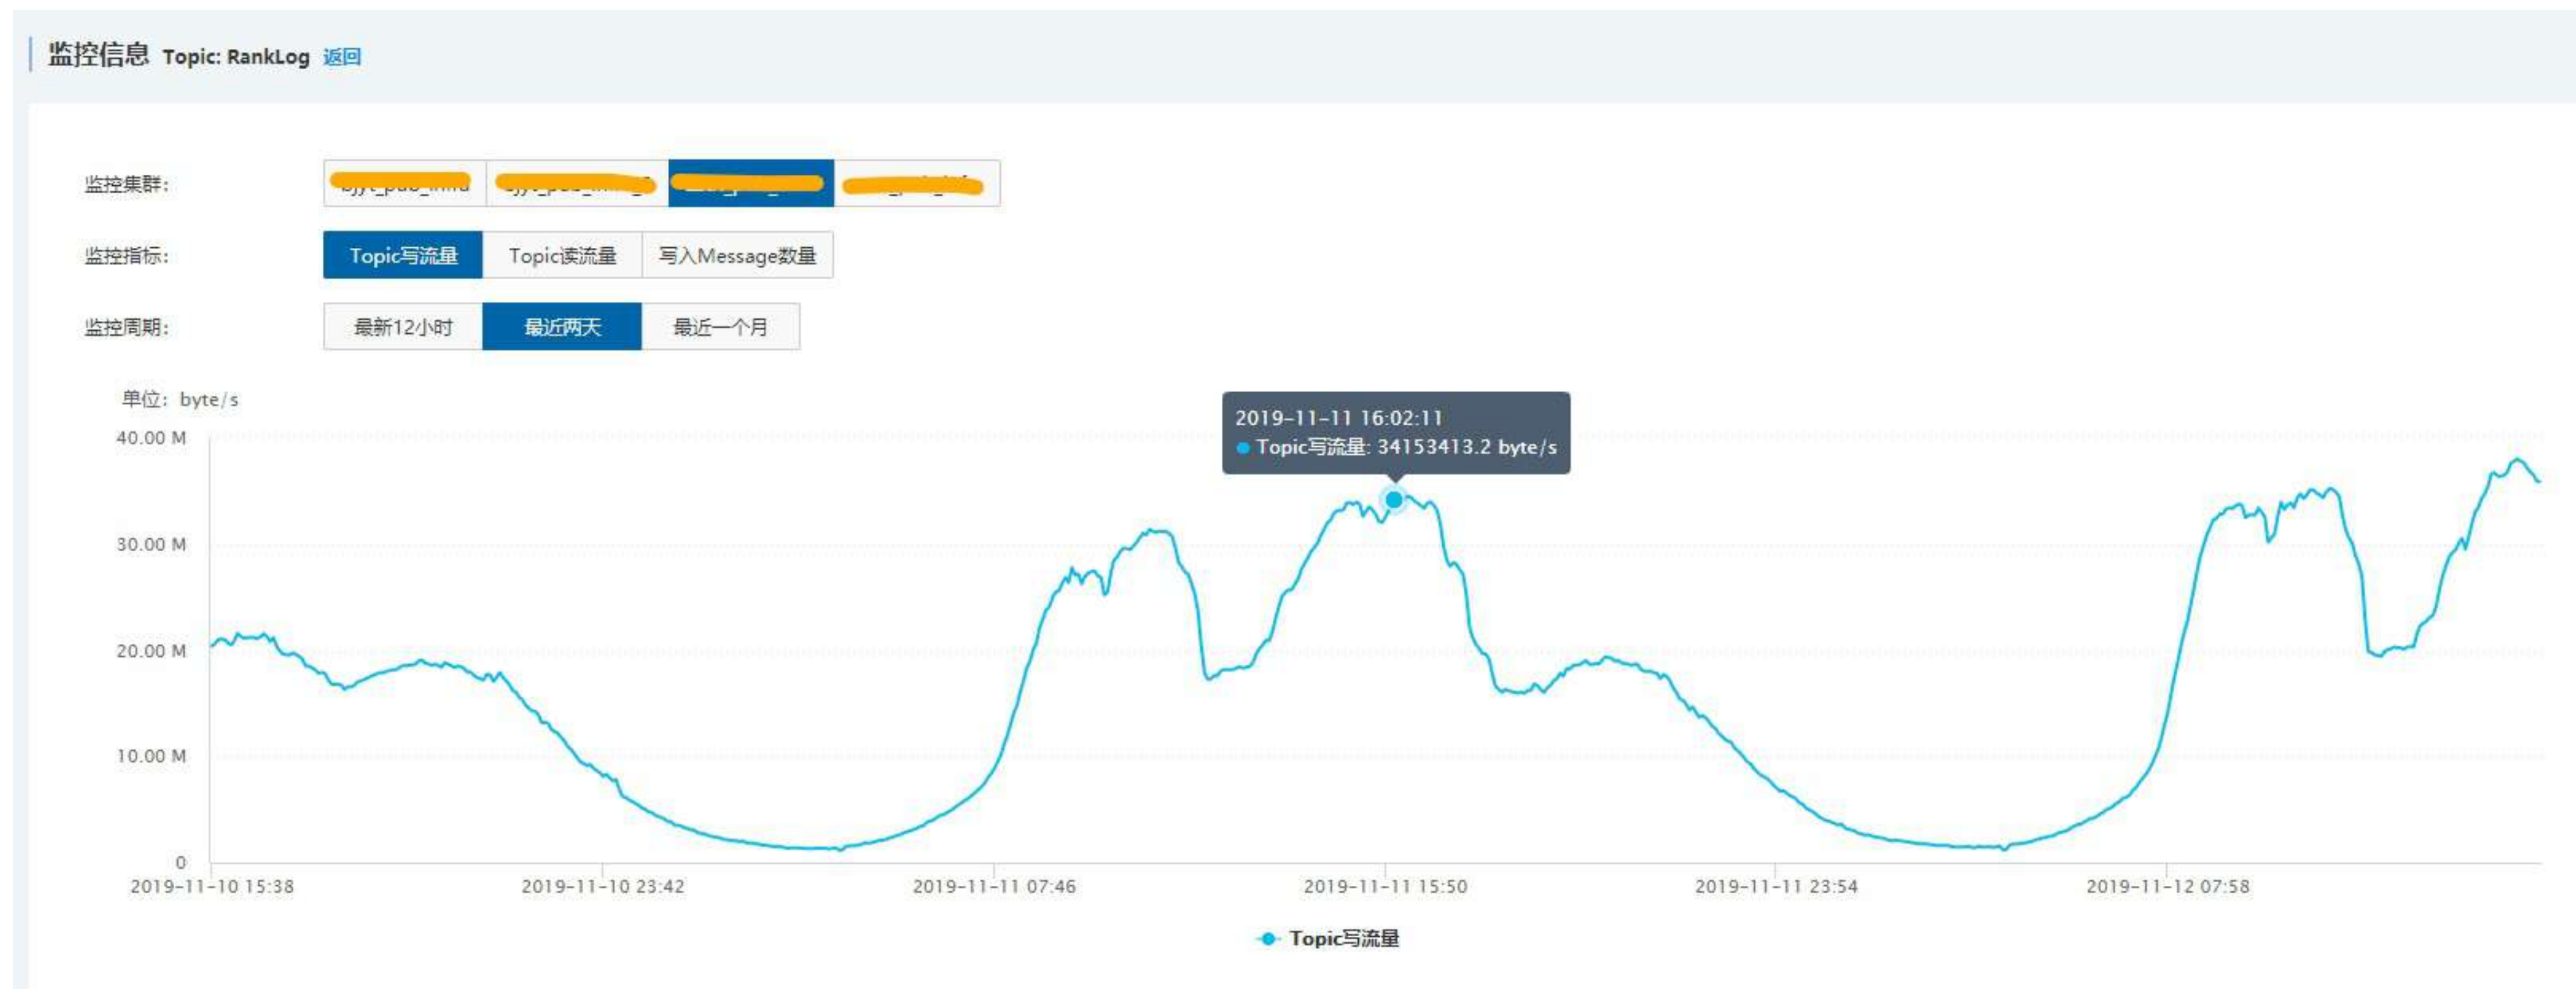
◆ 红红同学：通过调高并发总算是扛住了流量高峰，要不要把多余的资源还回去呢？

honghong: We managed to keep up with the traffic peak by increasing concurrency, do we need to return the extra resources?



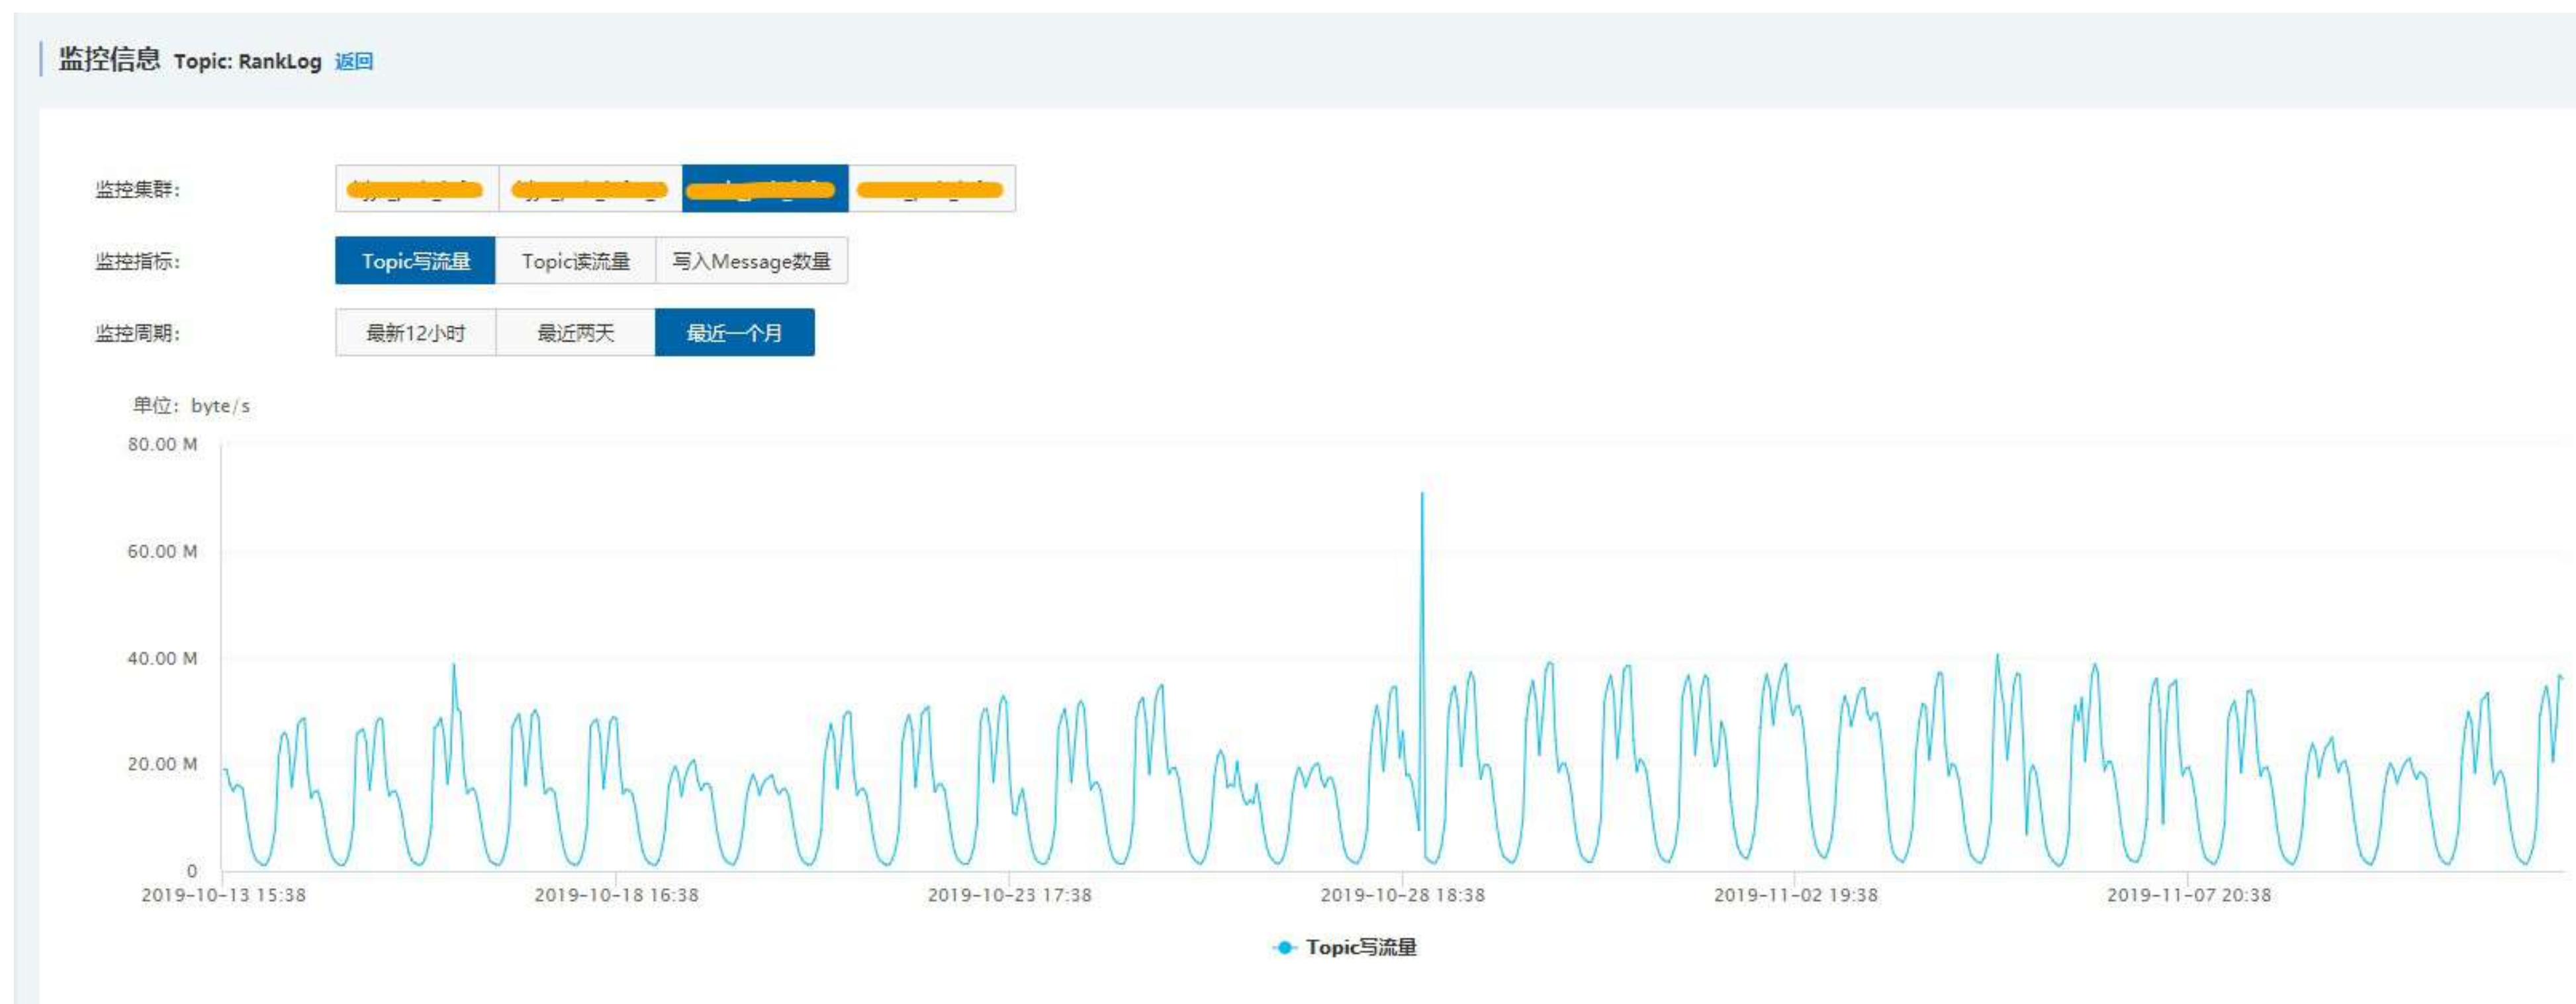
# 背景

## Background



# 背景

## Background





# 背景

Background

流式计算作业在未来都将不可避免的出现资源供应过剩或不足的情况！

Any streaming job will inevitably become over- or under-provisioned in the future.

何时调整， 如何调整？

When and how much to scale?



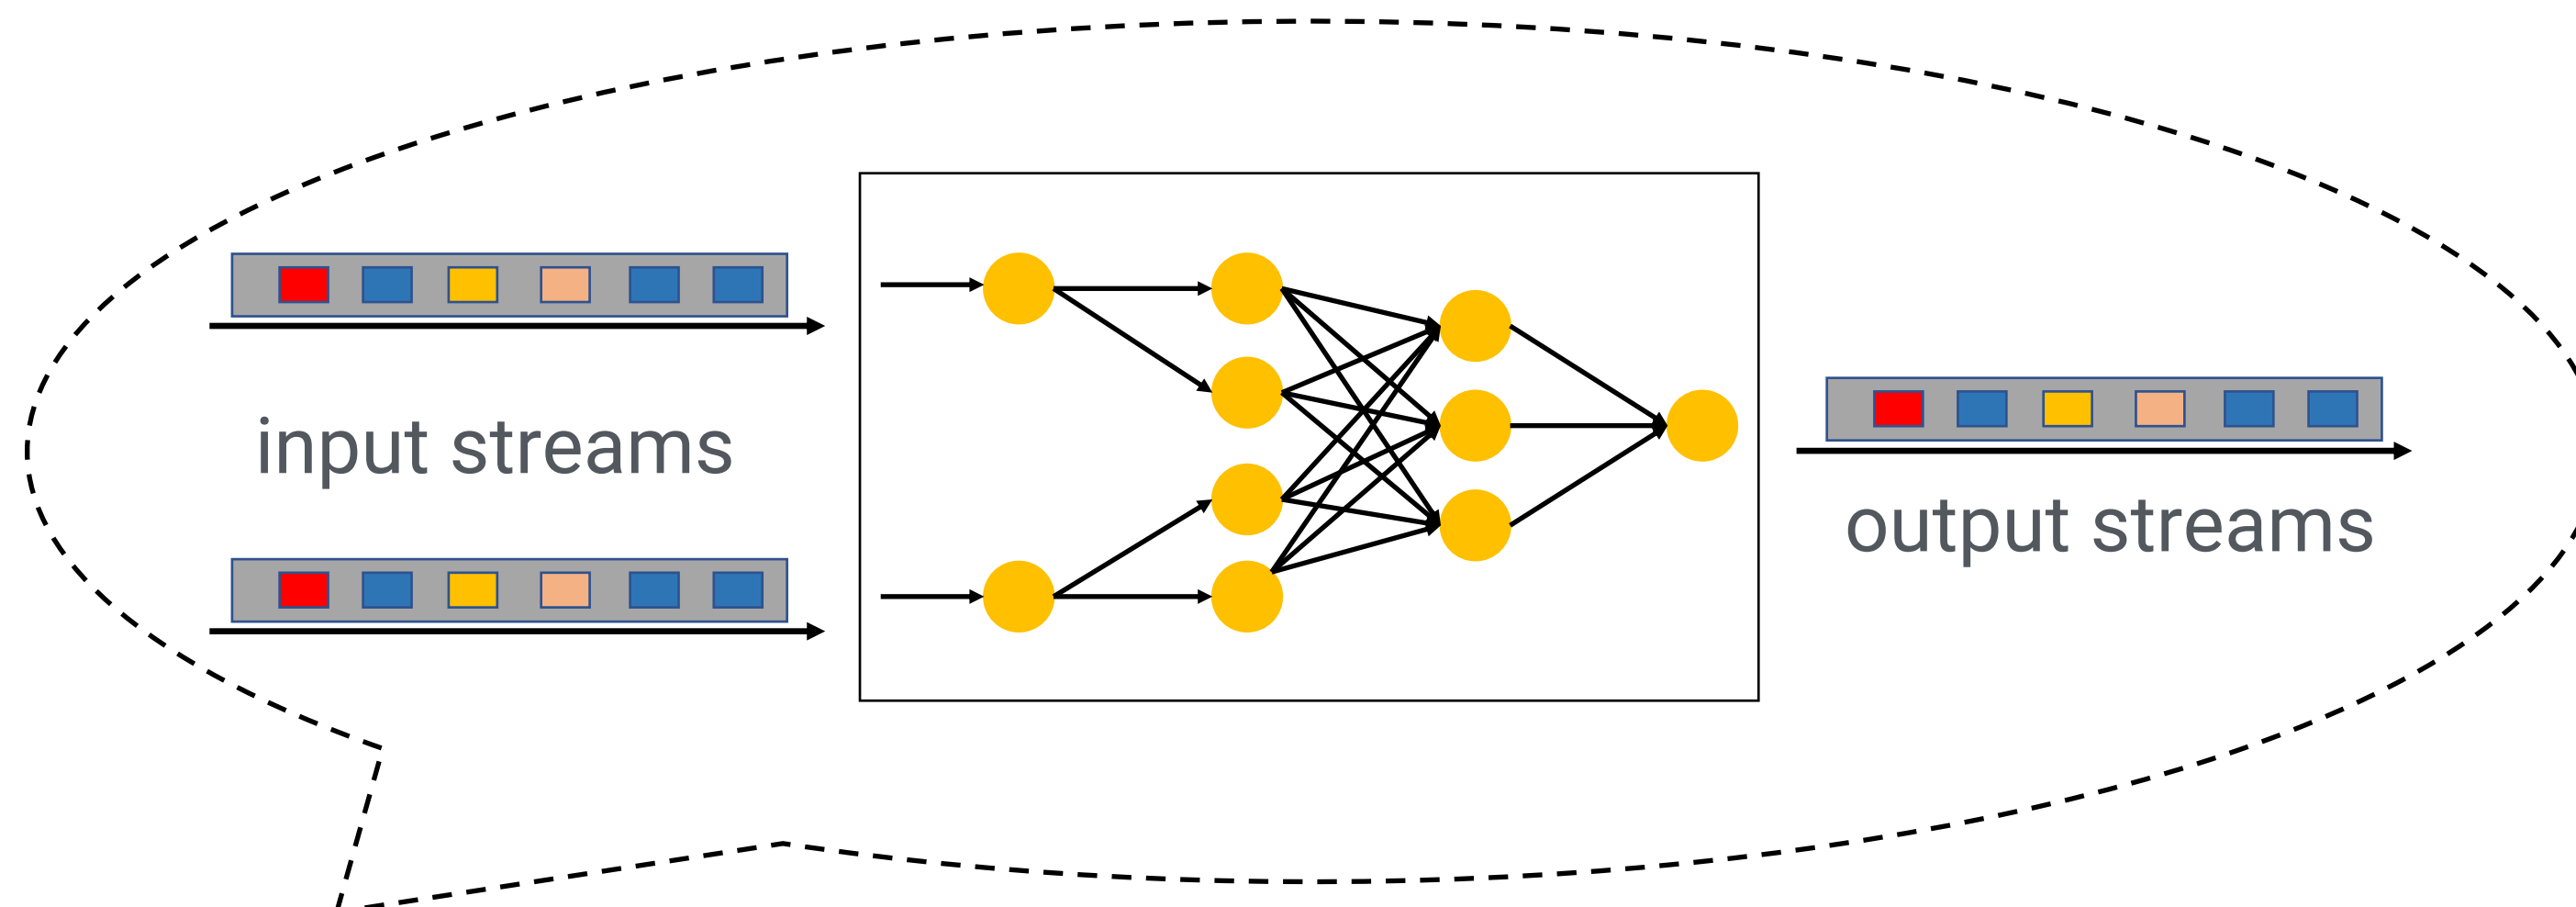
# 背景

Background

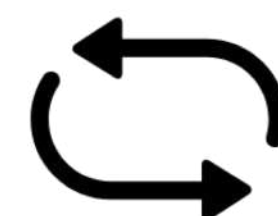
## 手动调整

Manual scaling overview

1. 监控源速率  
monitor event rates



2. 配置并行度  
configure parallelism



3. 部署并测试性能  
deploy and test performance

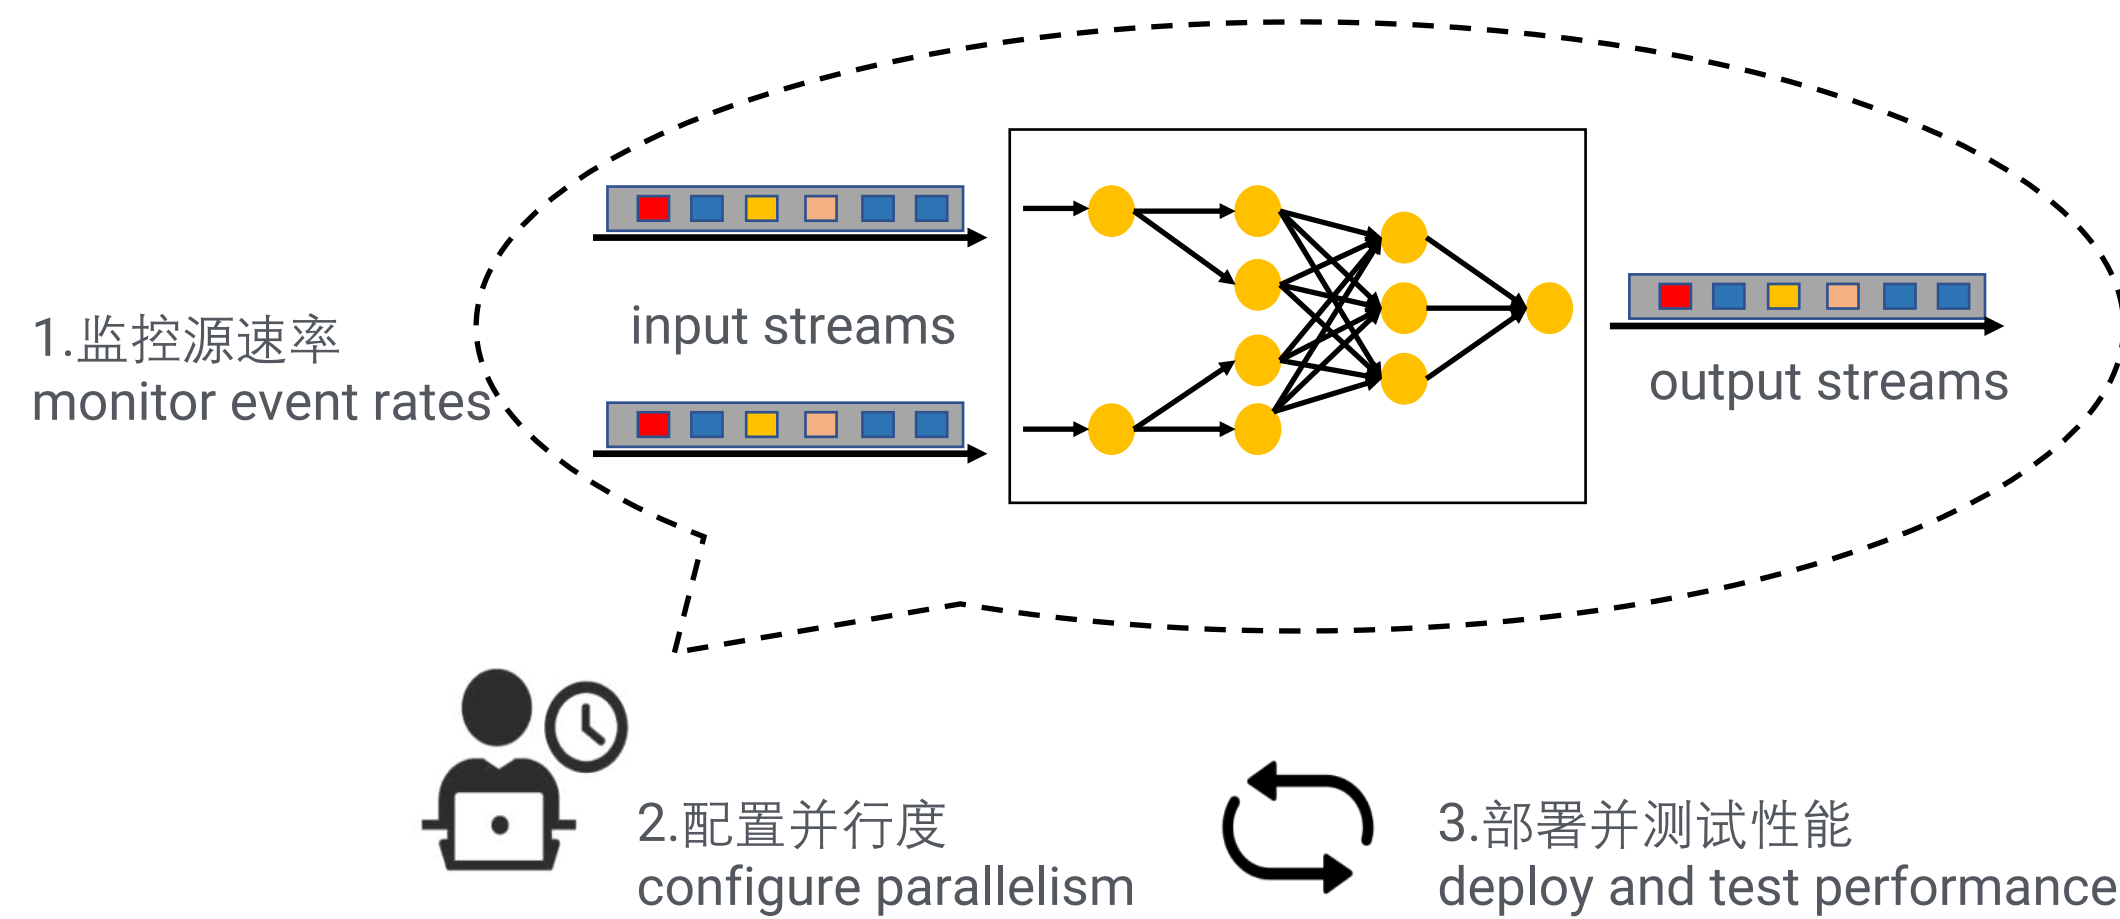


# 背景

Background

## 手动调整

Manual scaling overview



## 问题：

Problems:

- ◆ 额外的监控程序  
extra monitoring
- ◆ 凭经验，不准确  
depend on experience, inaccurate
- ◆ 影响作业稳定性  
affect job stability
- ◆ 手动配置难度大  
Manual configuration is difficult
- ◆ 手动调整成本高  
Manual scaling costs are high

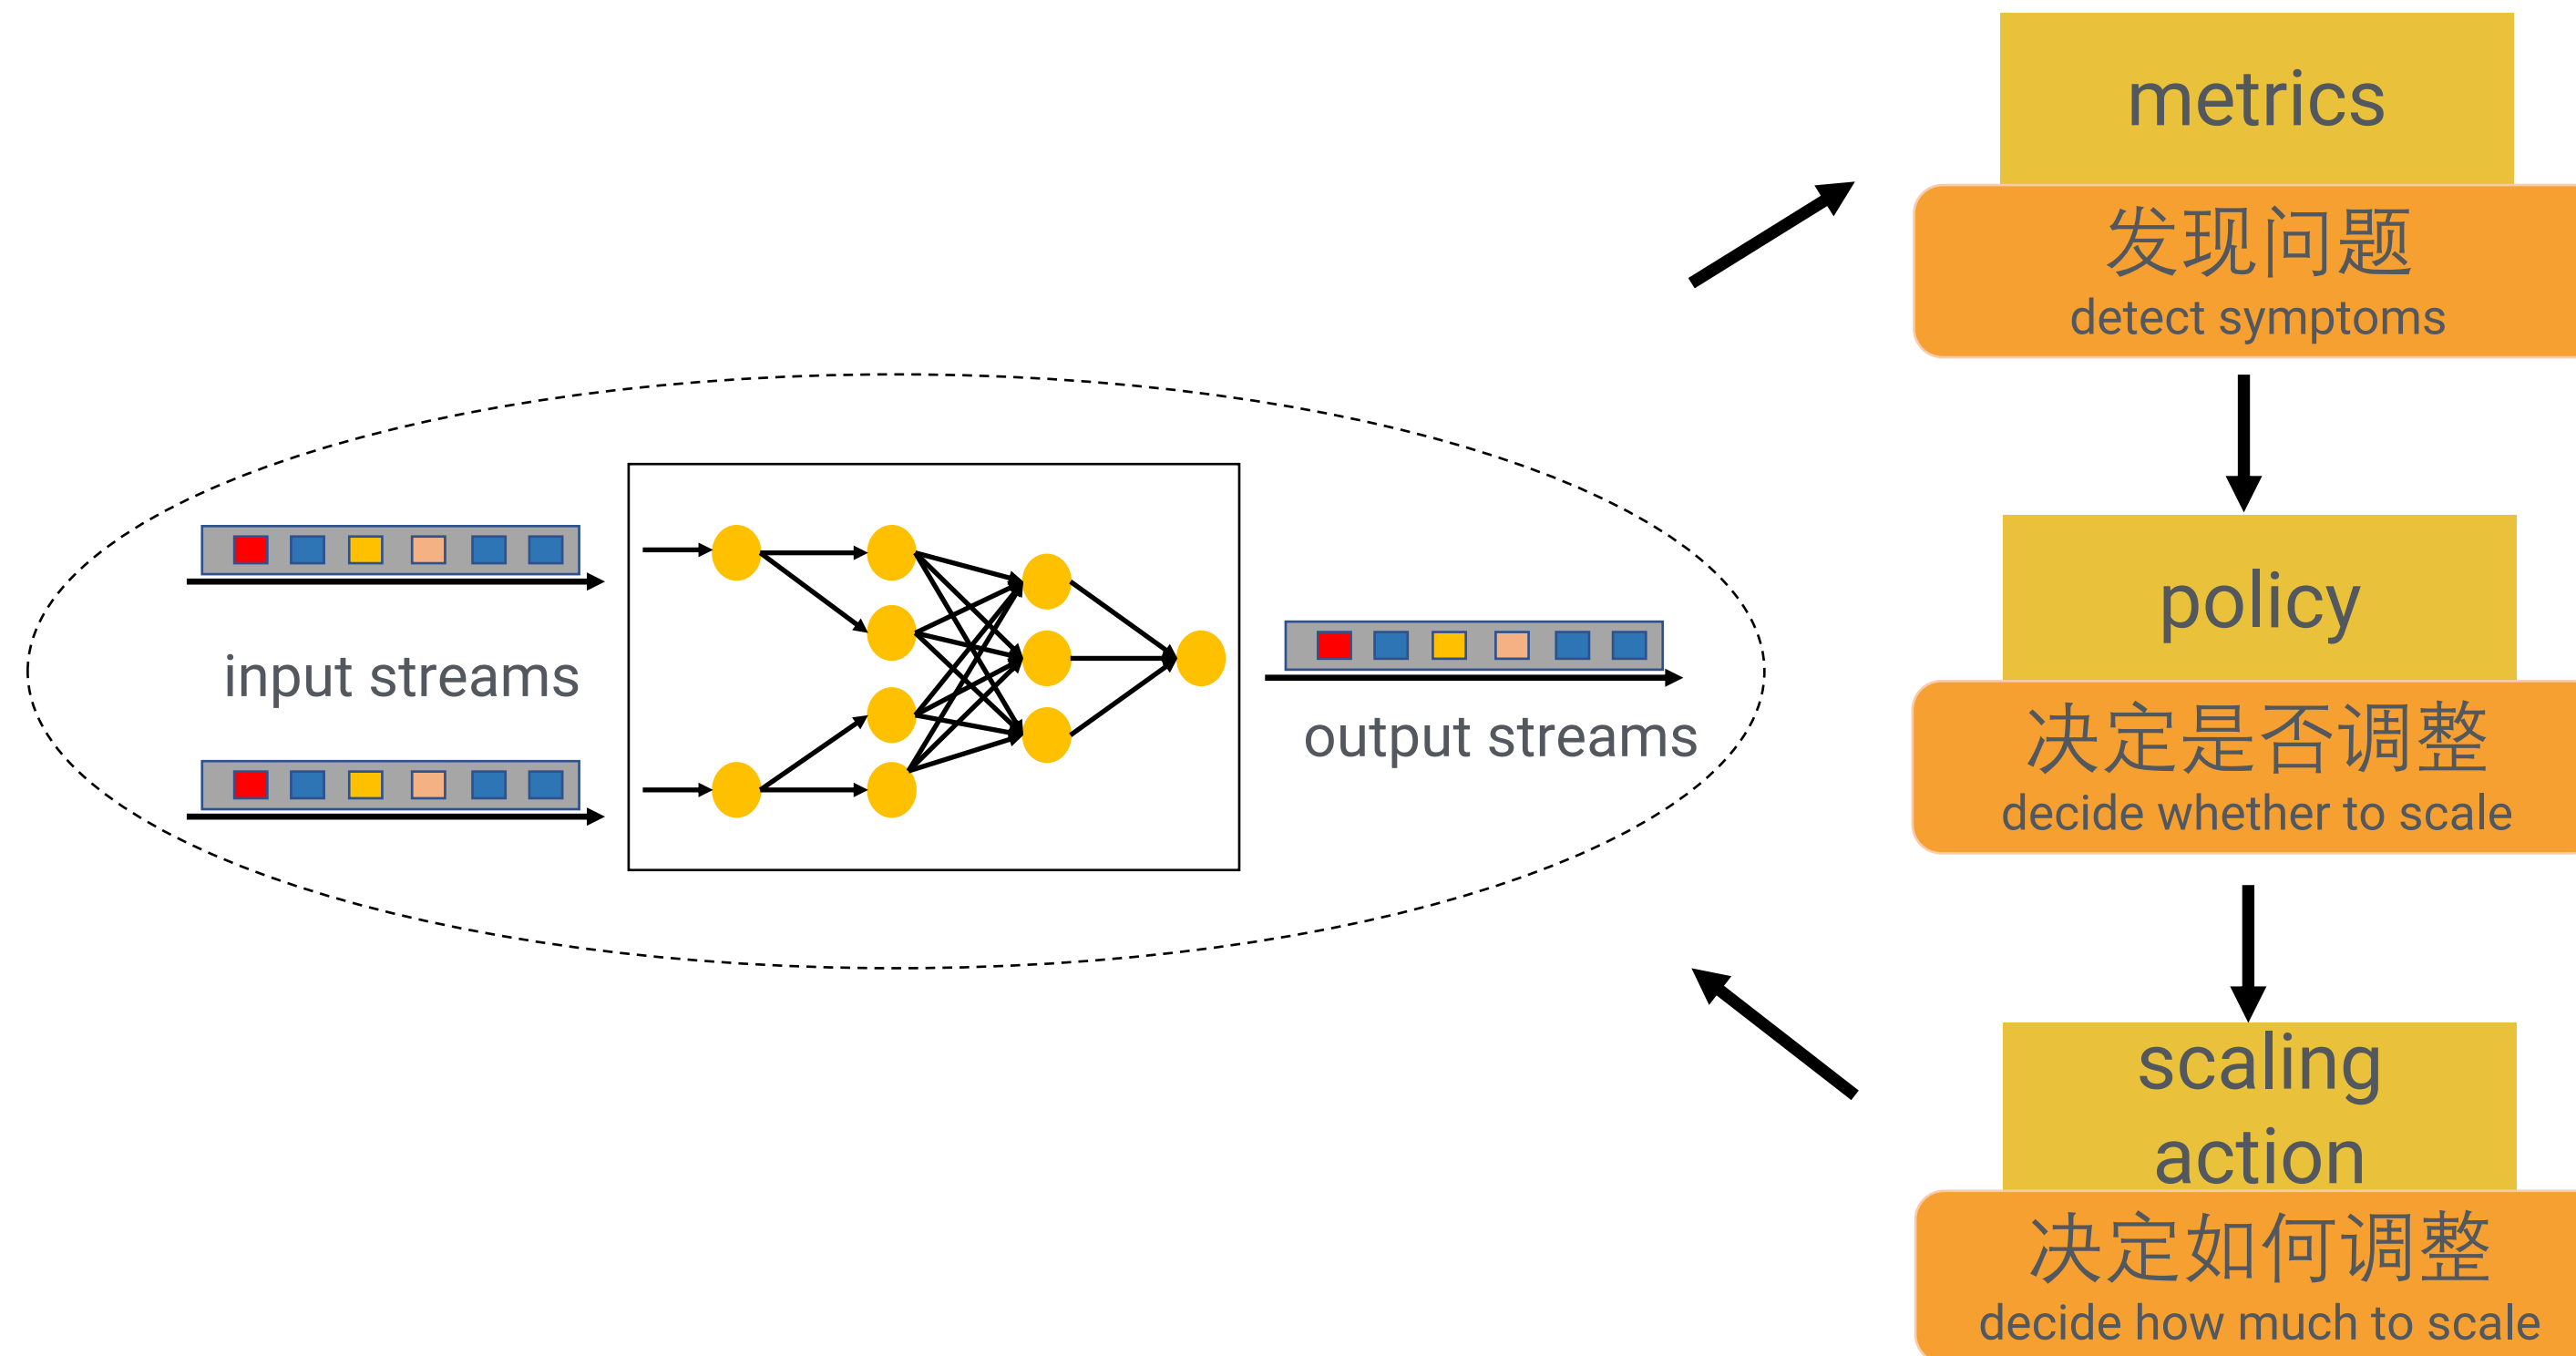


# 背景

Background

## 自动调整

Automatic scaling overview





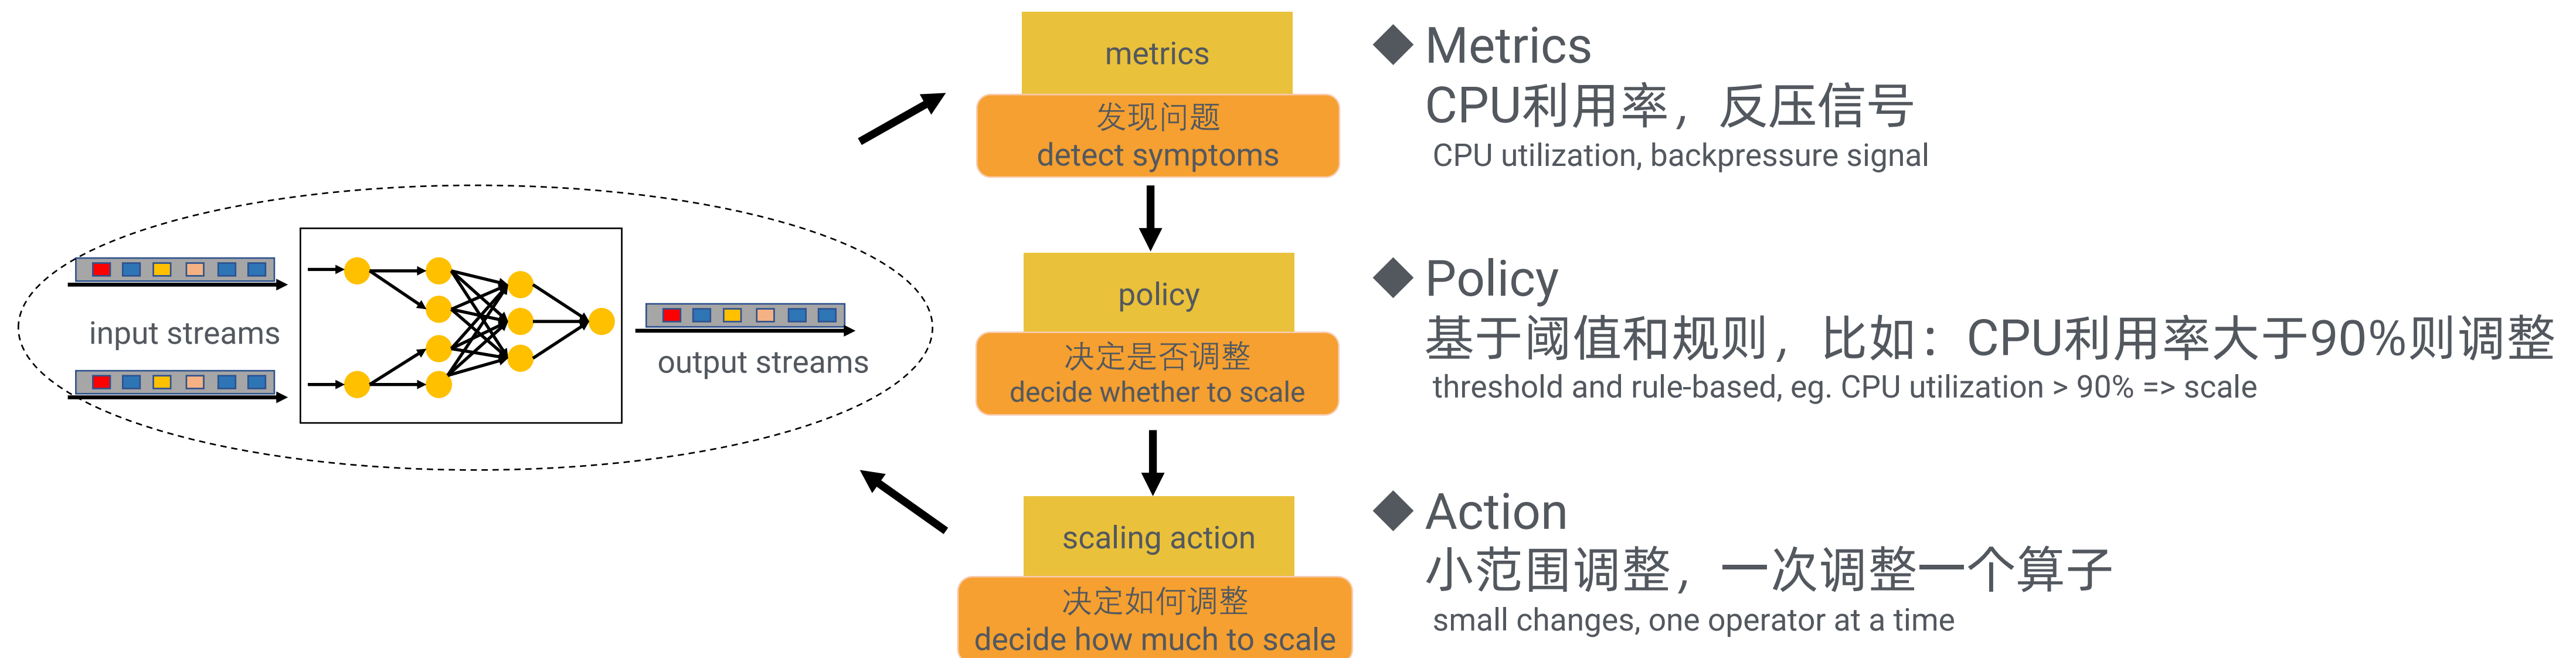
# 背景

Background

## 自动调整

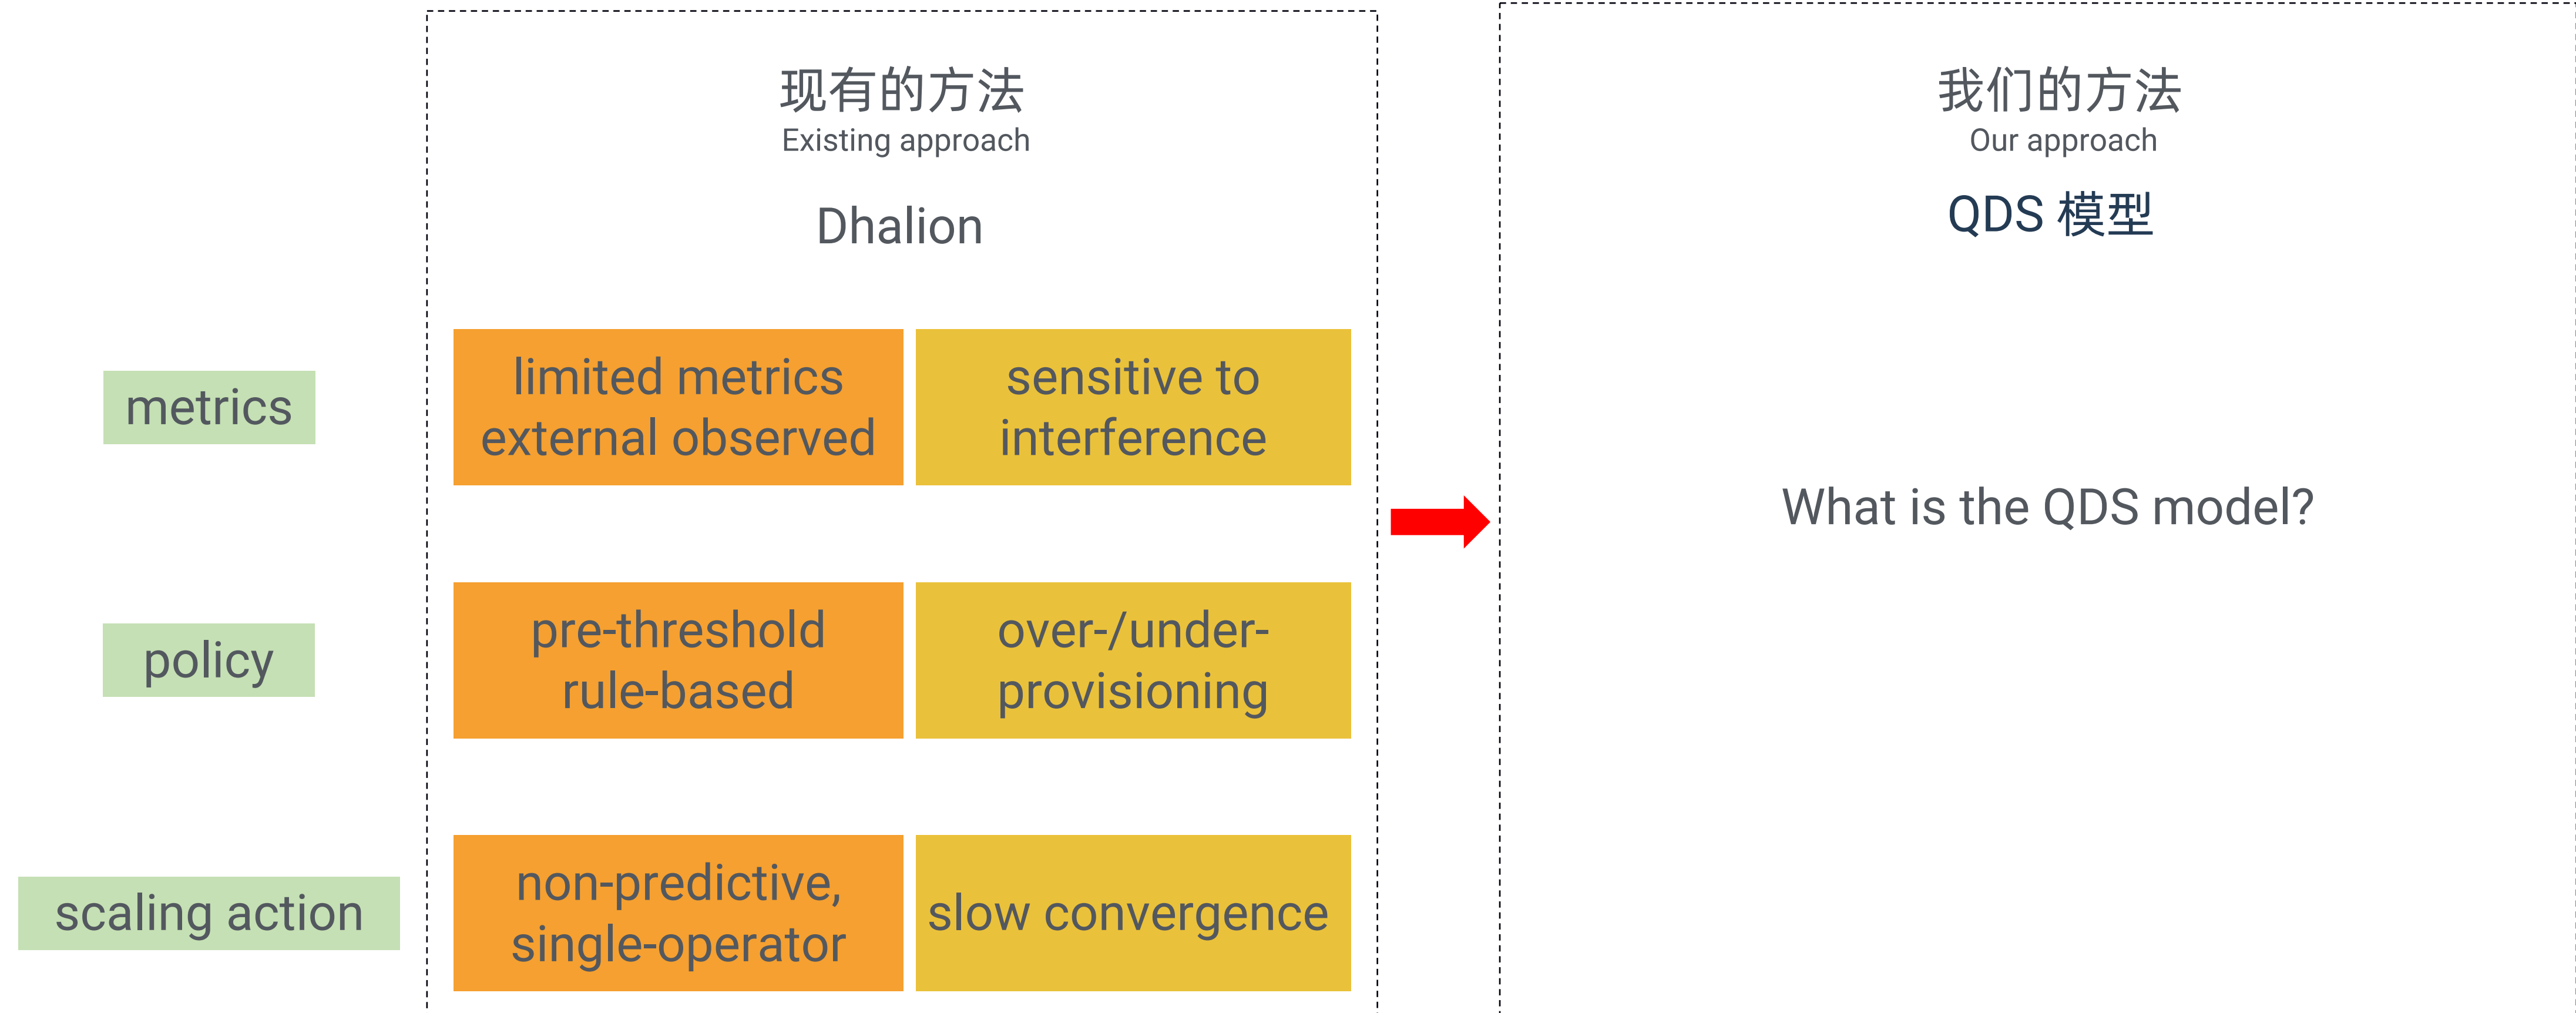
Automatic scaling overview

Dhalion



# 背景

Background





# QDS 模型

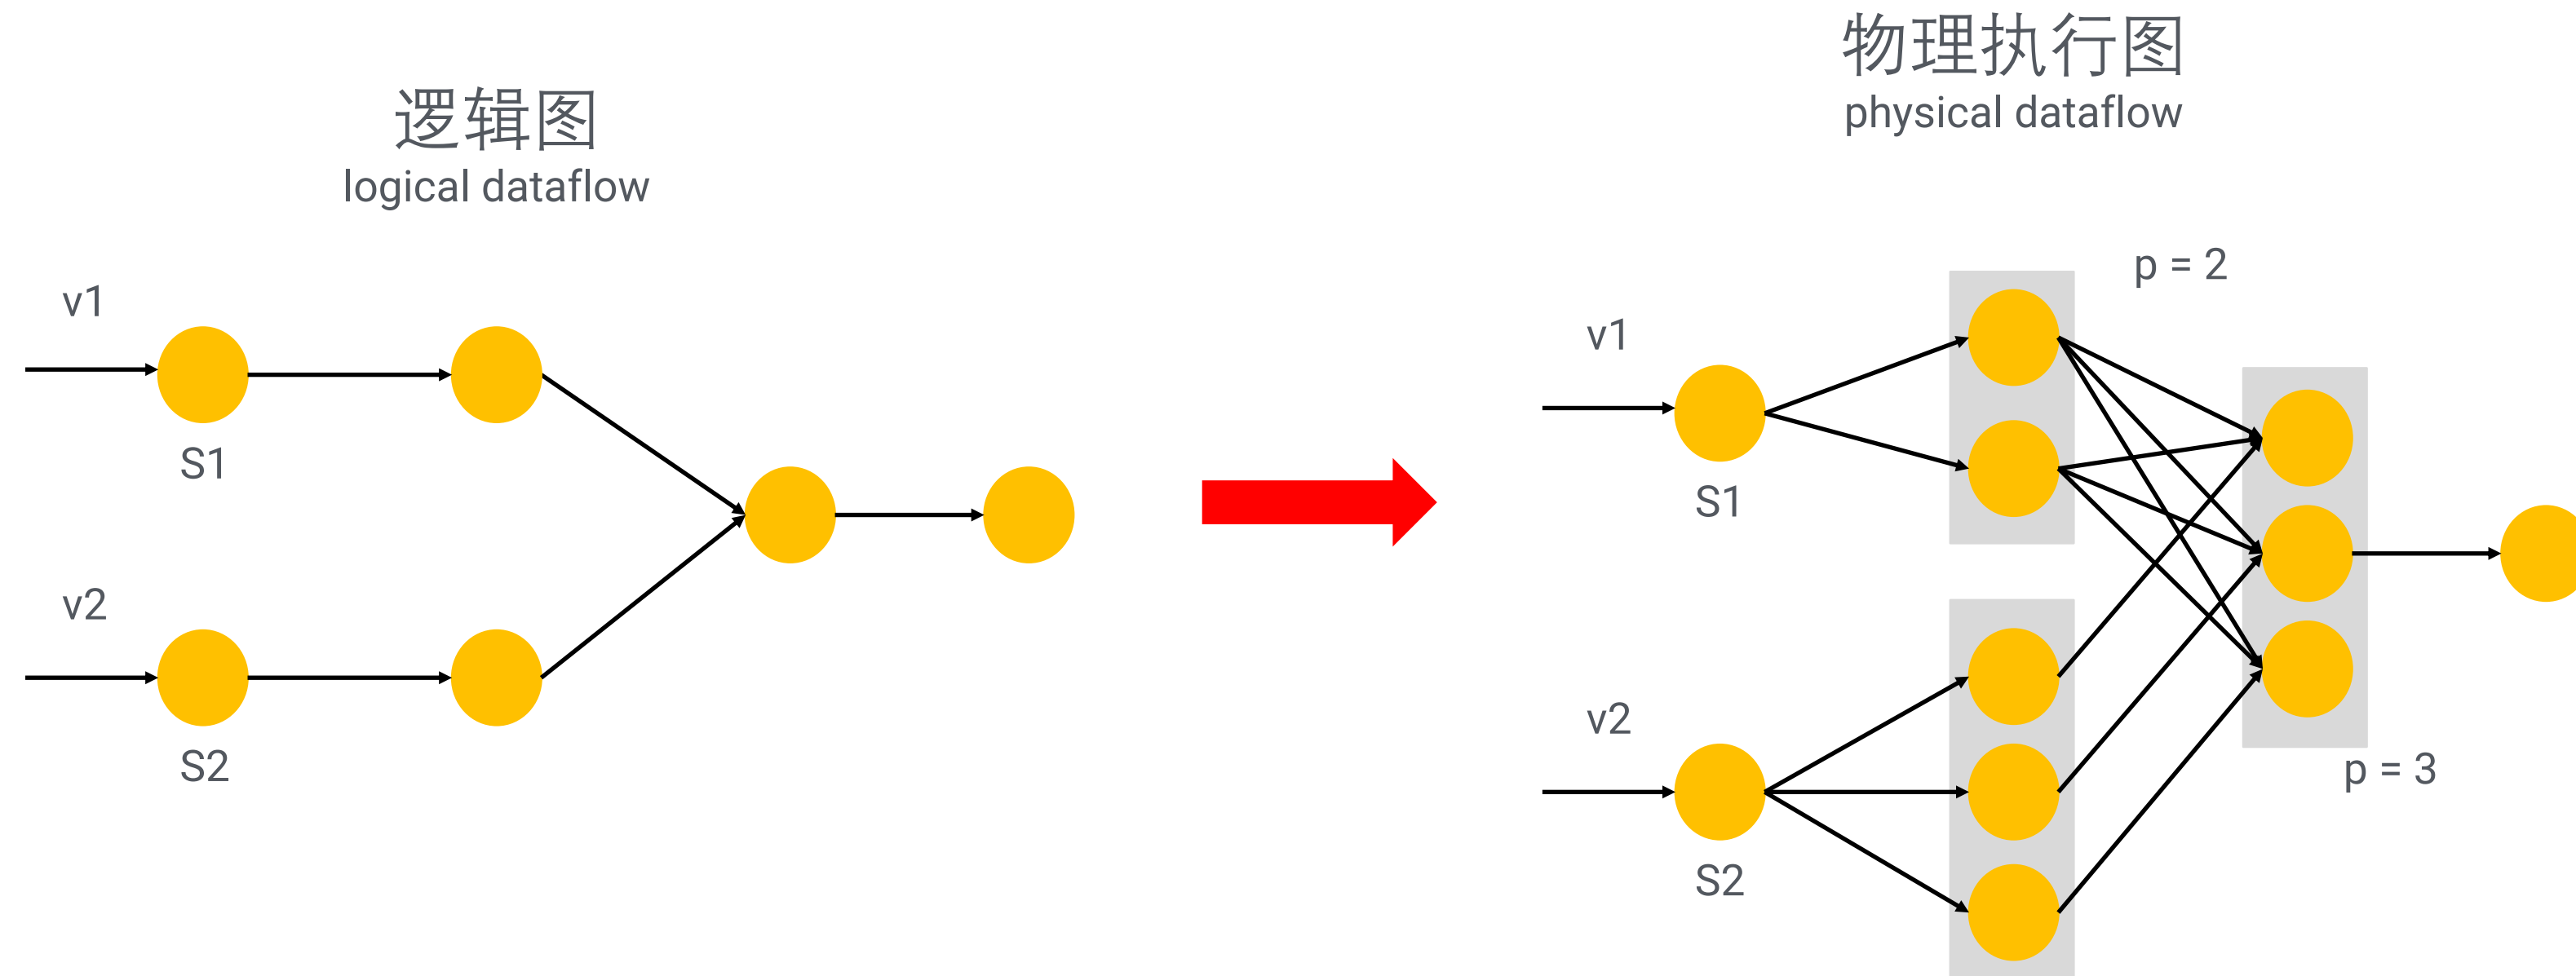
The QDS model

---

02

# 目标

Target



目标 1：在保证维持源速率的情况下确定各个算子的最小并行度！

Target 1: identify the minimum parallelism per operator such that the physical dataflow can sustain all source rates!

目标 2：在保证源没有堆积的情况下确定各个算子的最小并行度！

Target 2: identify the minimum parallelism per operator such that the physical dataflow have no data backlogs!



# 介绍

## Introduction

### 我们的方法

Our approach

### QDS 模型

true rate through  
instrumentation

no oscillations

dataflow  
dependency model

no  
over-/under-shoot

predictive,  
dataflow-wide action

fast convergence

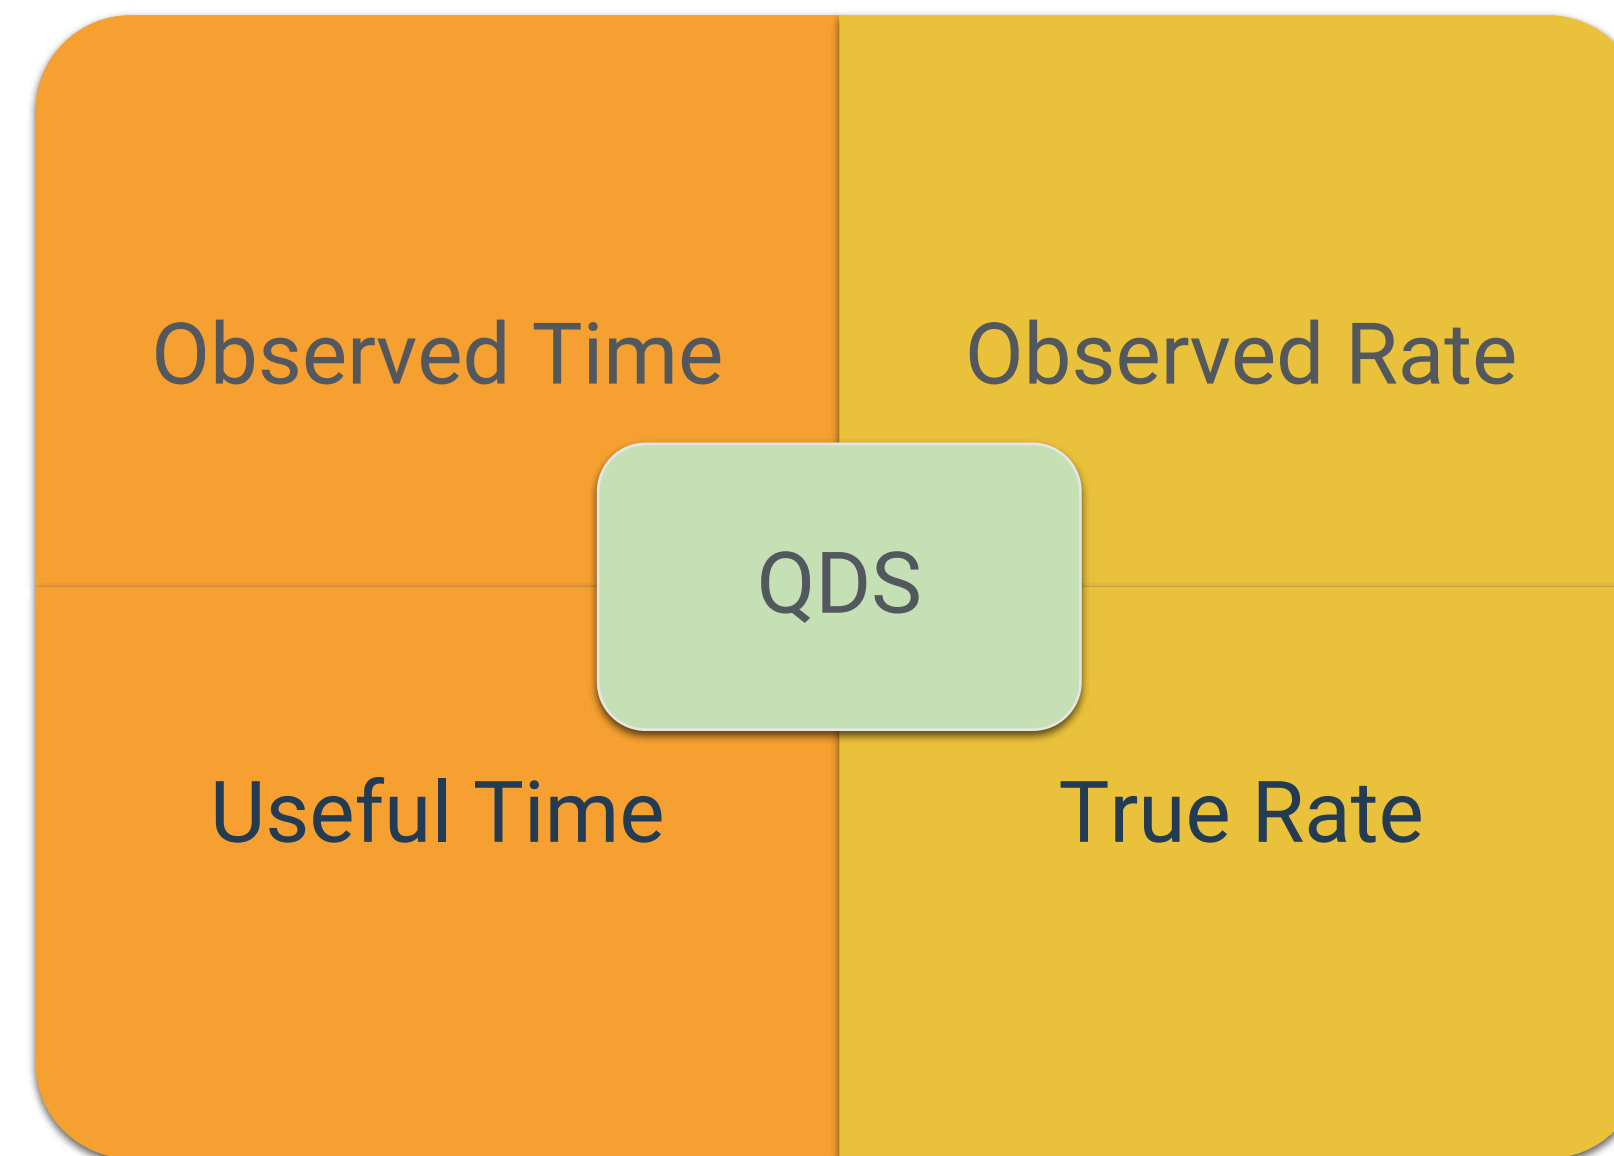
➤ 关注数据源  
consider the sources

➤ 关注每个算子的真正的处理能力和输出能力  
consider each operator's true processing and output capabilities

➤ 关注数据流本身，各个算子间的计算依赖性  
consider the dataflow topology and computational dependencies among operators

# 核心概念

The core concept

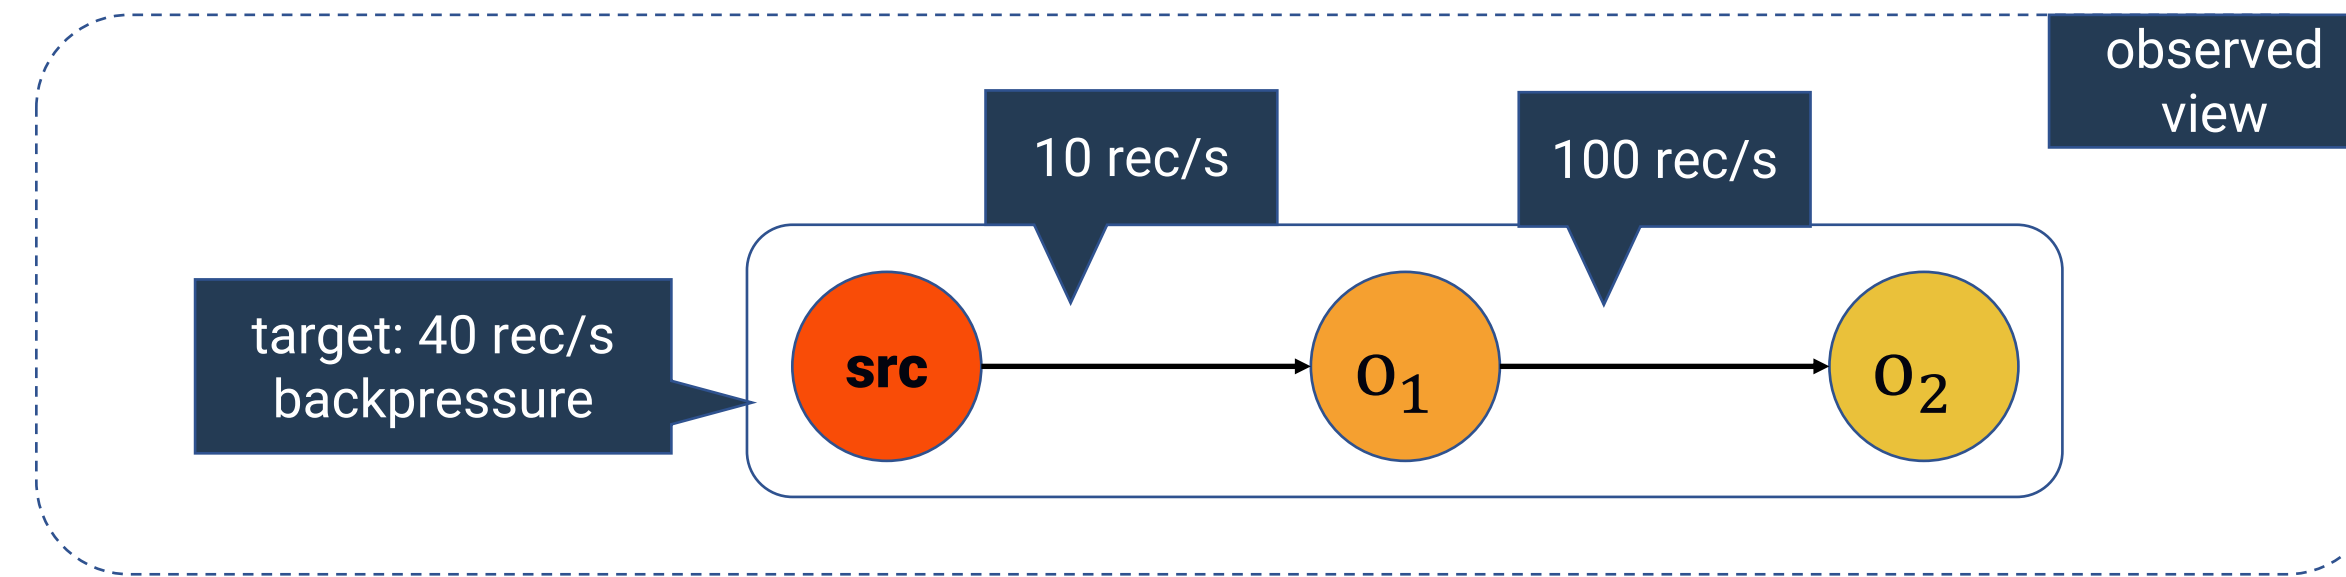


$$\text{Optimal parallelism for } o_i = \frac{\text{aggregated true output rate of upstream ops}}{\text{average true processing rate of } o_i}$$



# 举个例子

For example



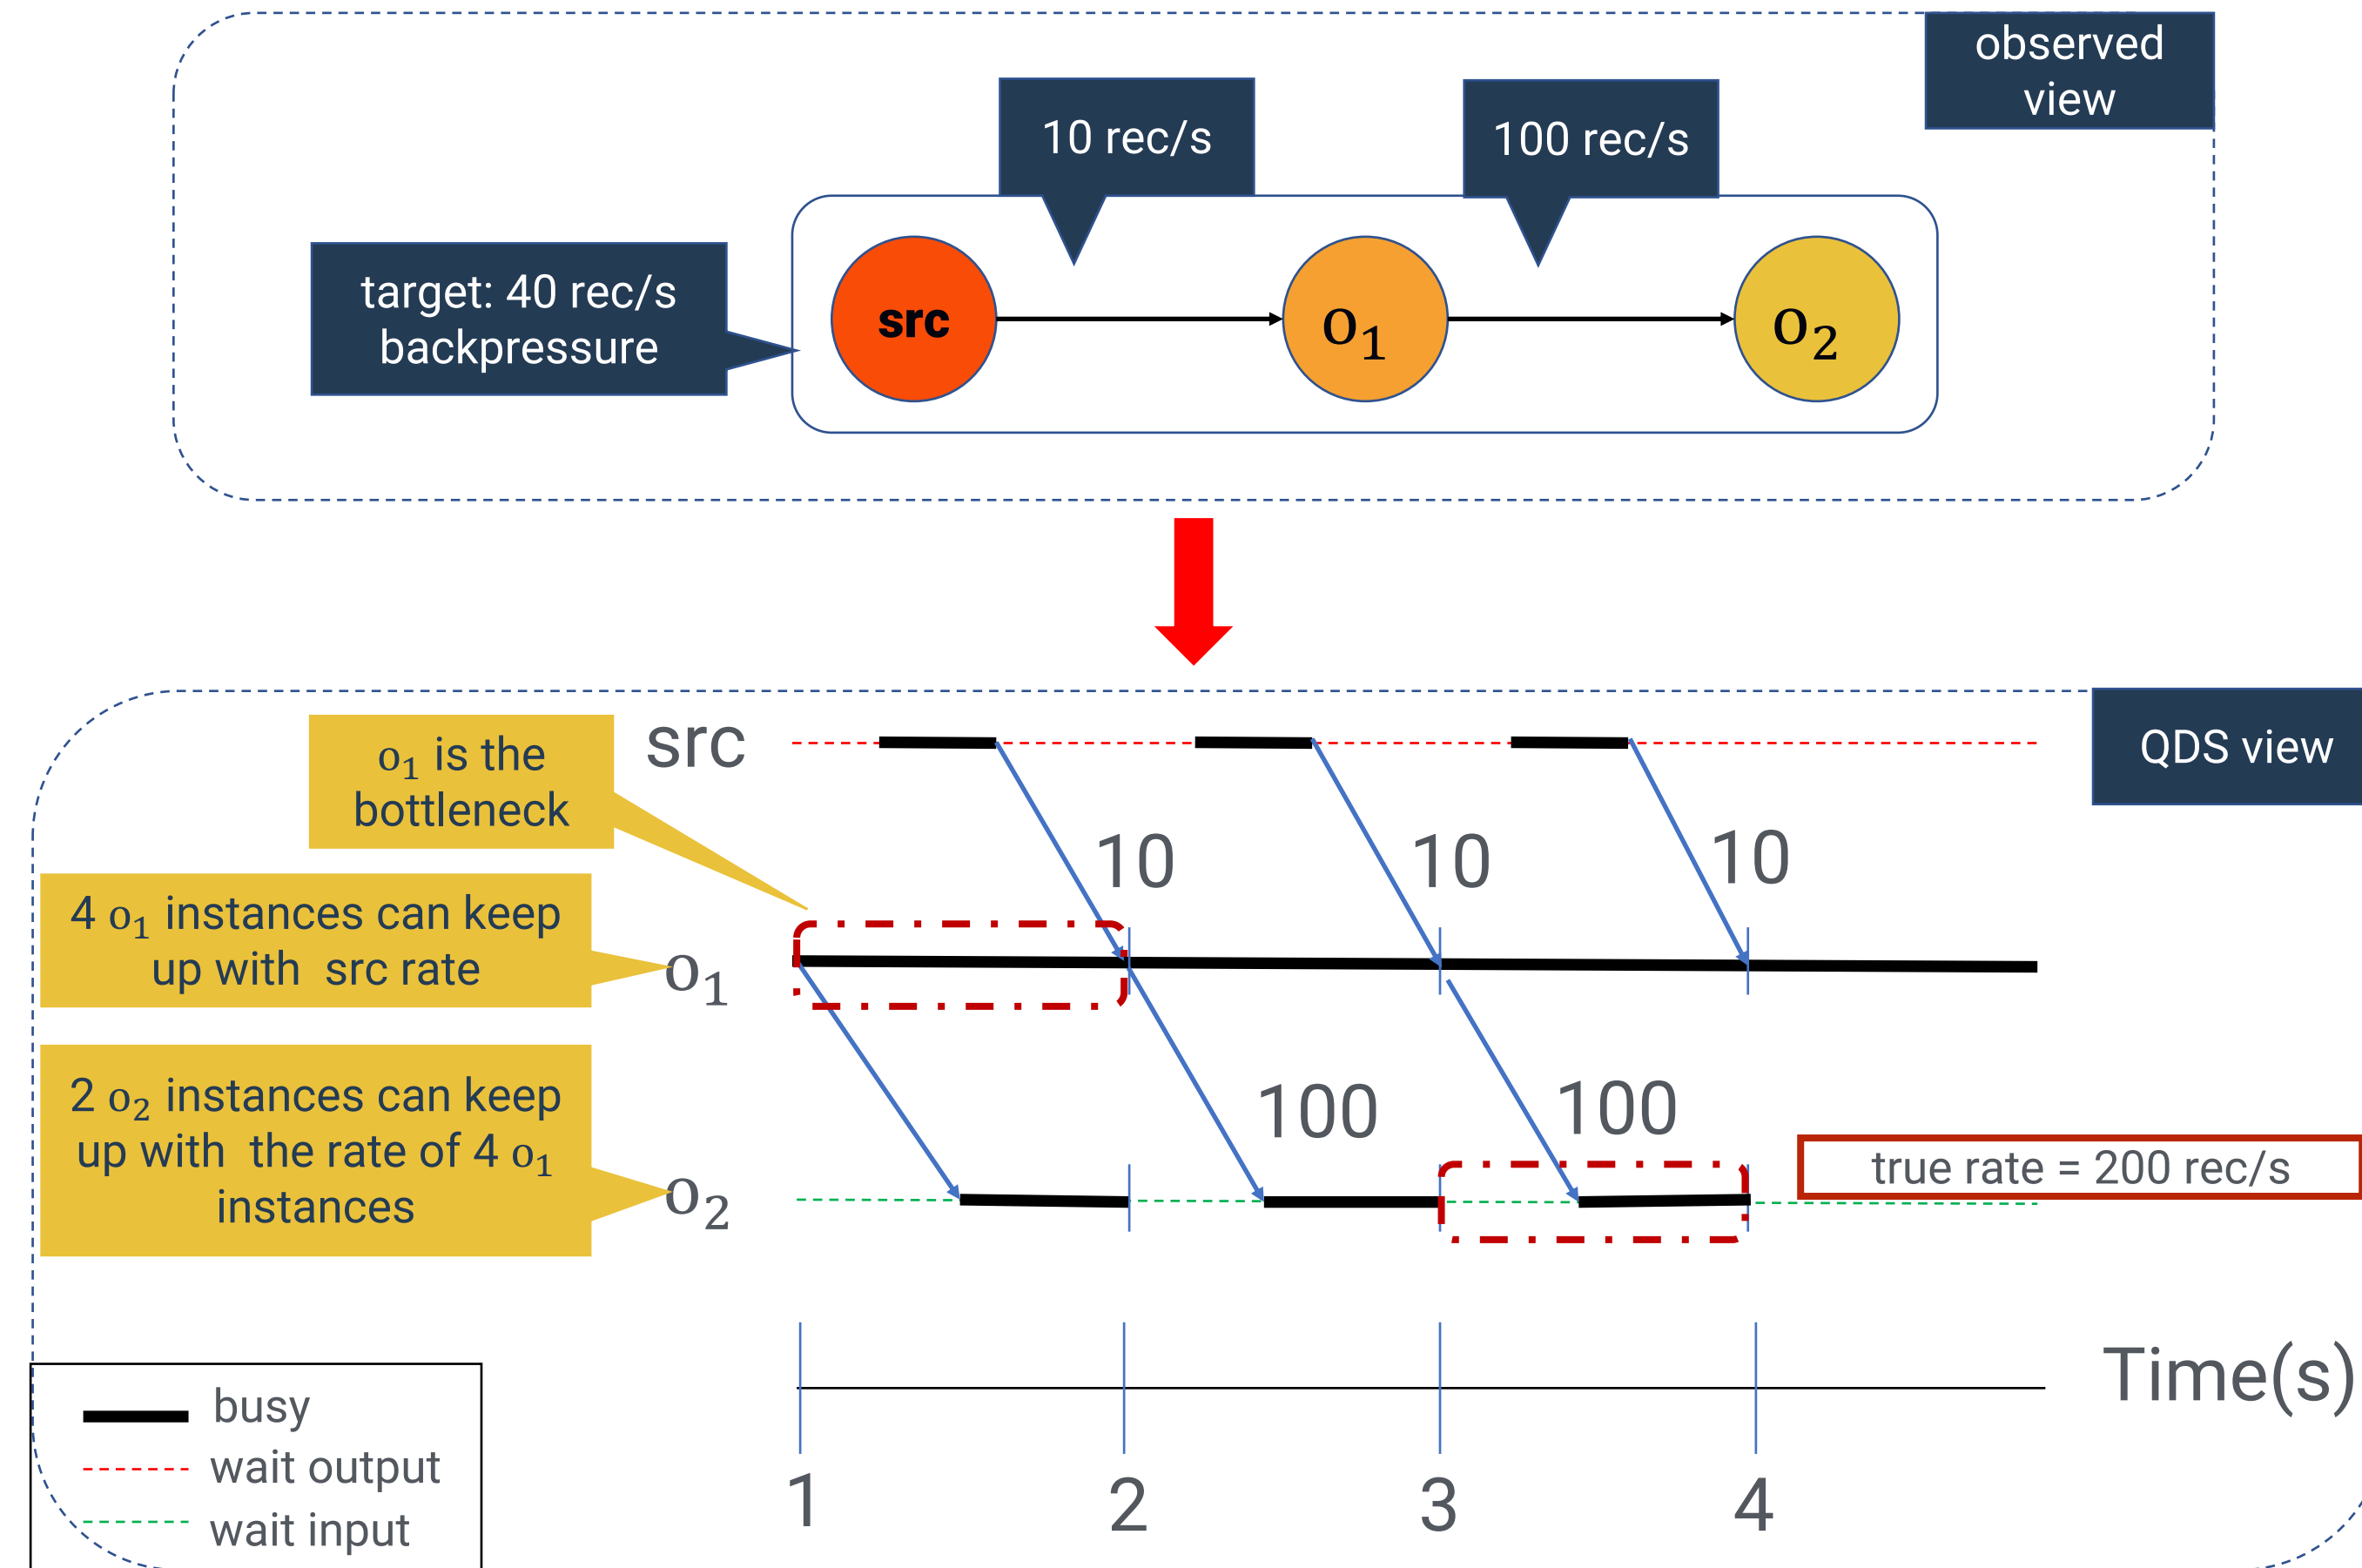
Which operator is the bottleneck?

What if we scale o<sub>1</sub> x 4?

How much to scale o<sub>2</sub>?

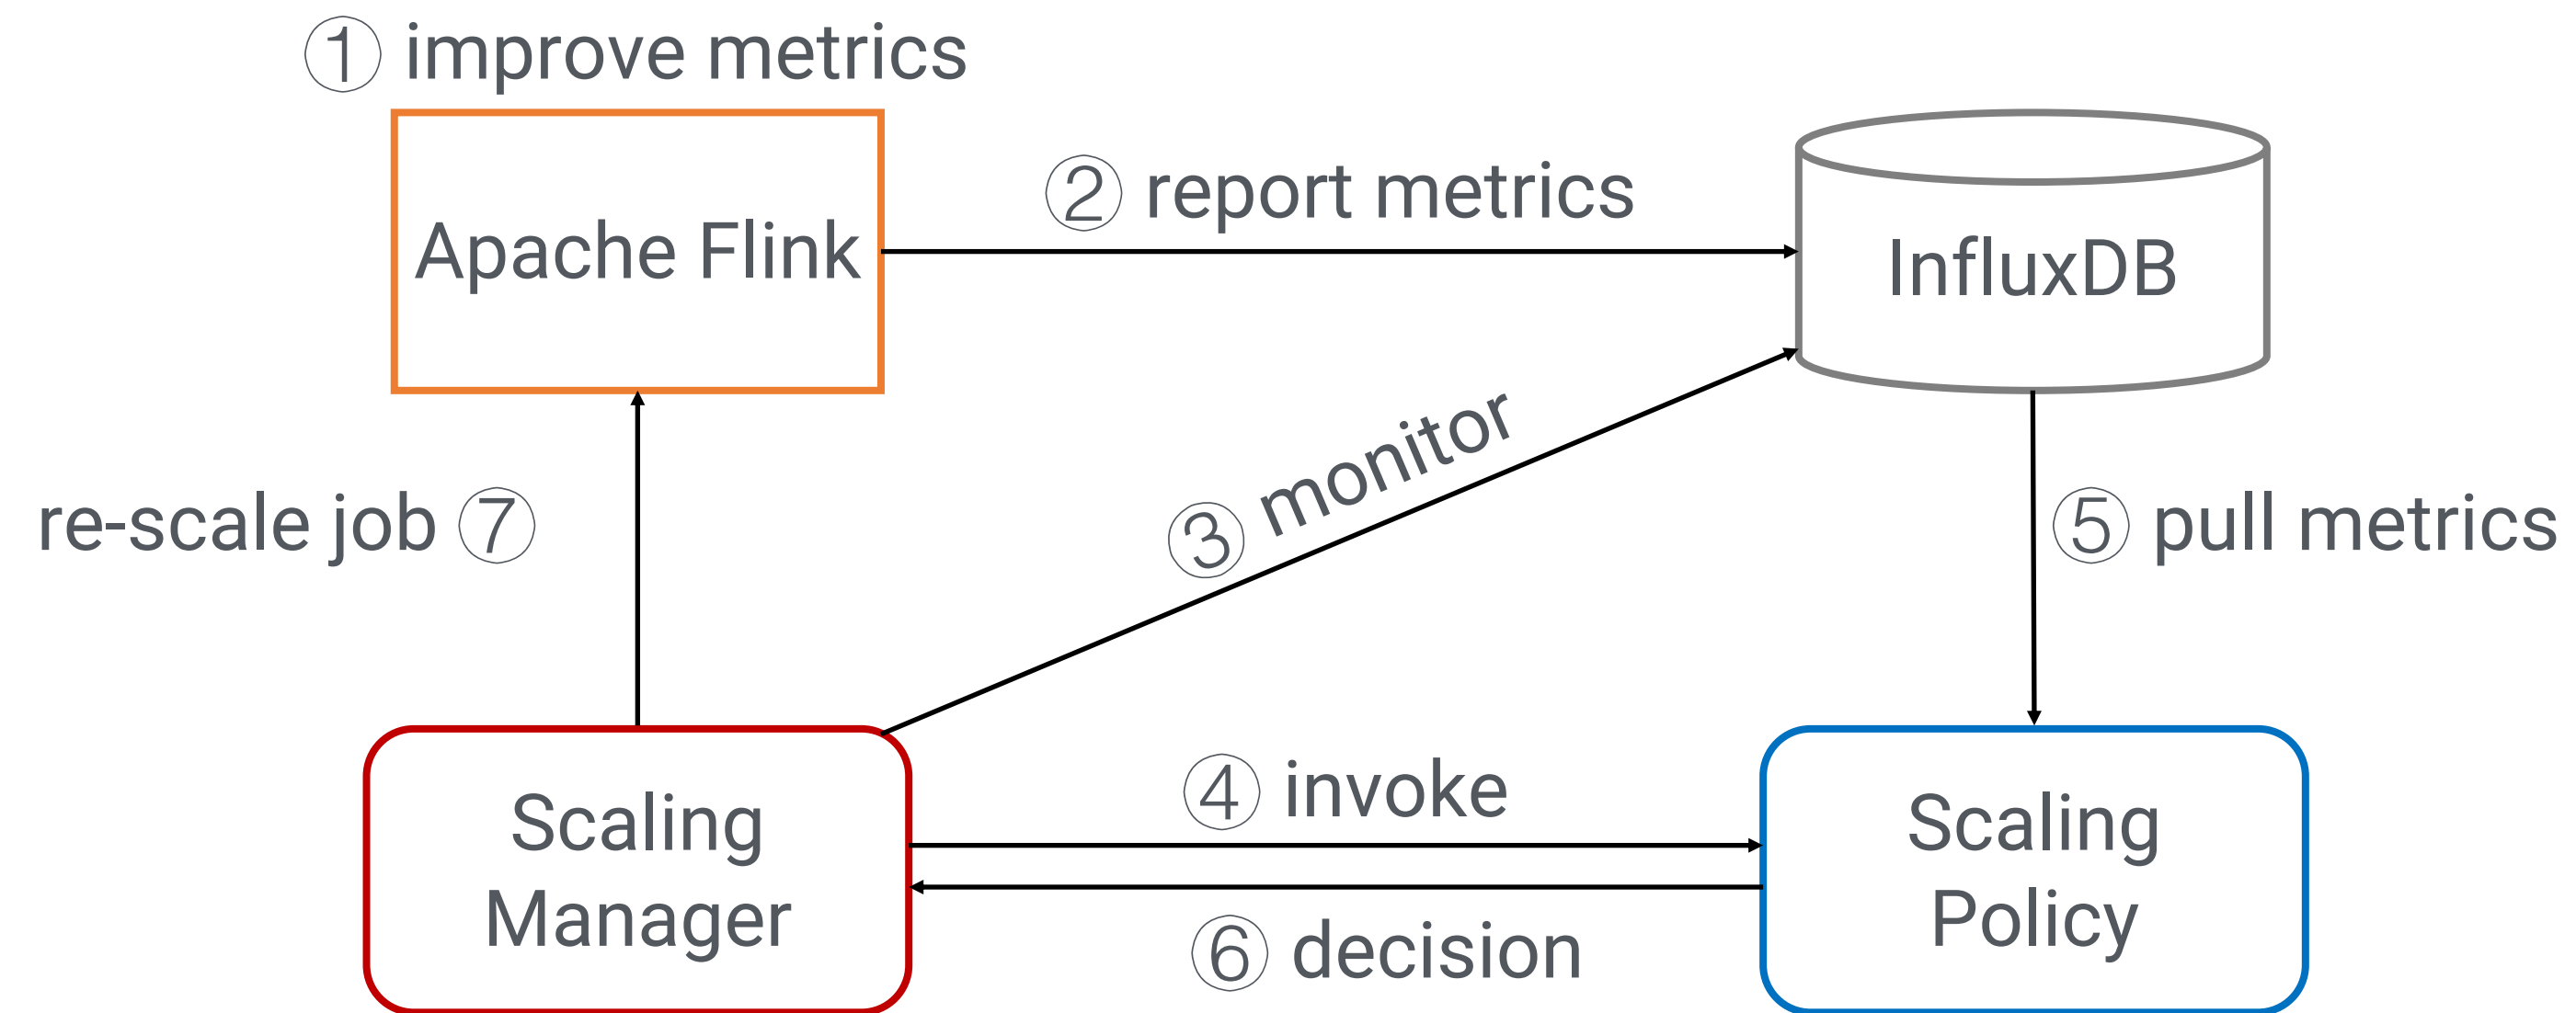
# 举个例子

For example





# QDS on Flink



# 生产实践

Production Practice

---

03



# 生产实践

## Production Practices

### 准备工作：

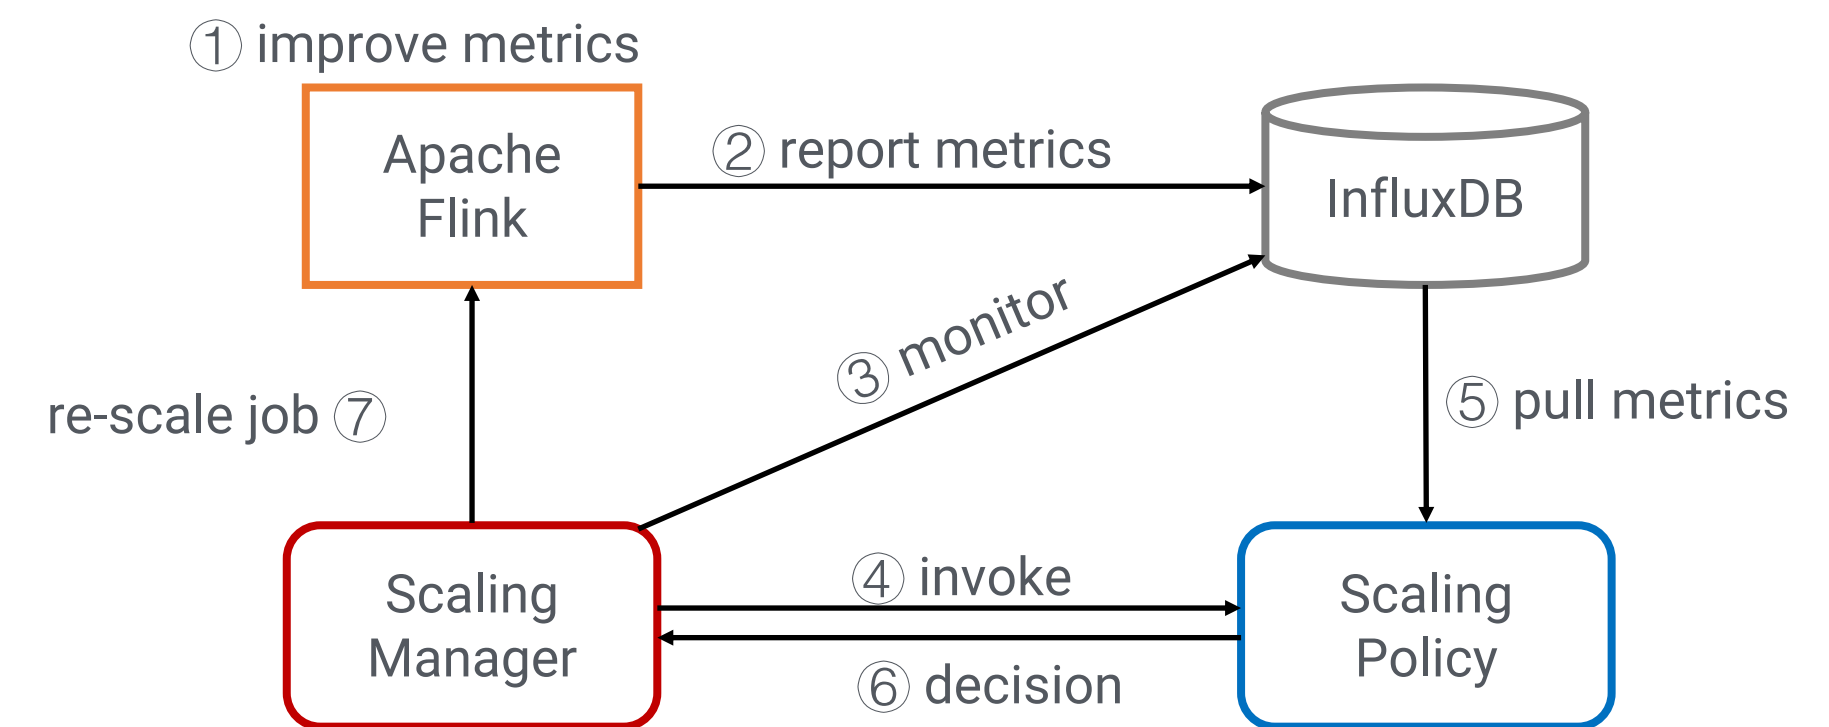
The preparatory work:

### 打开自动调整配置

Enable automatic scaling configuration

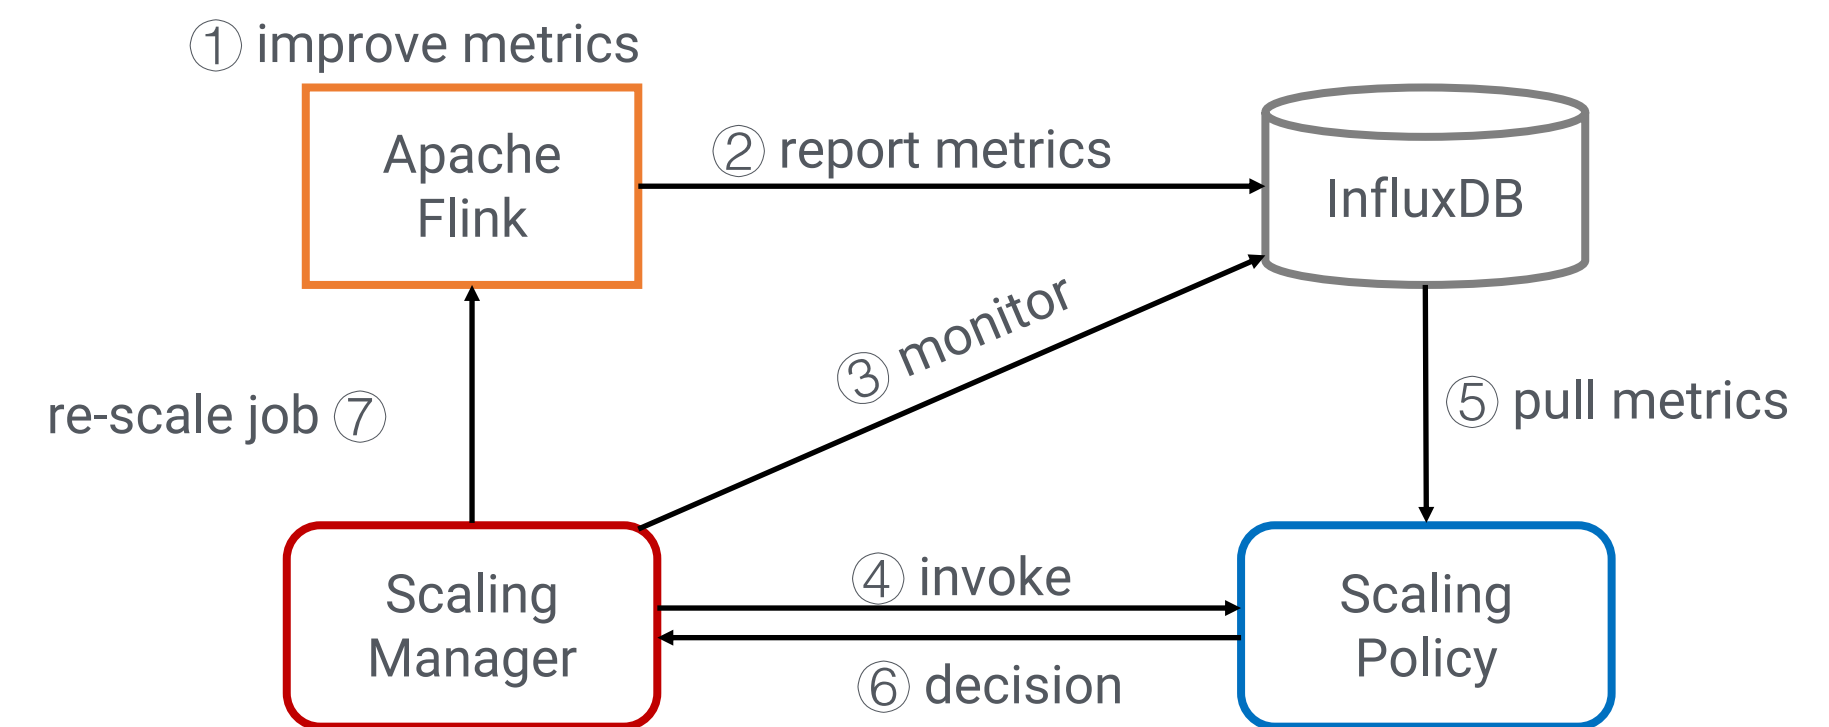
首次提交作业，各算子默认并行度为 1，不开启 chain 策略。

Each operator default parallelism is 1 in the first submission of the job, and disable chain.



# 生产实践

## Production Practices

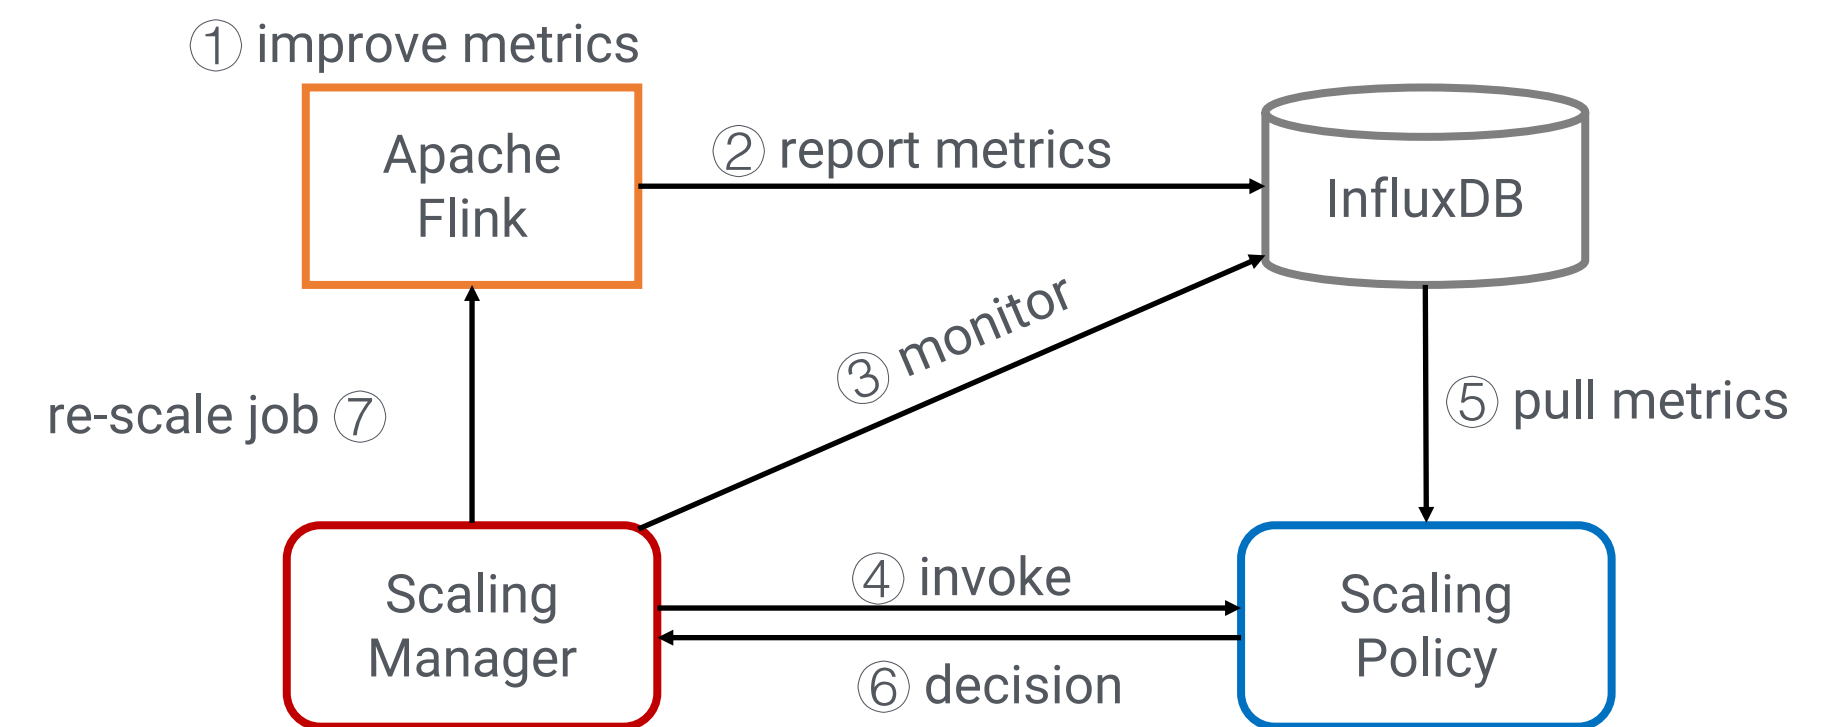


### ① improve metrics

- ✓ taskmanager\_job\_task\_operator\_KafkaConsumer\_topic\_partition\_0\_committedOffsets
- ✓ taskmanager\_job\_task\_operator\_KafkaConsumer\_topic\_partition\_0\_latestOffsets
  
- ✓ taskmanager\_job\_task\_operator\_deserializationDuration
- ✓ taskmanager\_job\_task\_operator\_processingDuration
- ✓ taskmanager\_job\_task\_operator\_serializationDuration
- ✓ taskmanager\_job\_task\_operator\_waitingDuration

# 生产实践

## Production Practices



### ② report metrics

✓ `org.apache.flink.metrics.influxdb.InfluxdbReporter`

将 `/opt/flink-metrics-influxdb-1.9.0.jar` 拷贝到 `/lib` 目录下  
Copy `/opt/flink-metrics-influxdb-1.9.0.jar` into the `/lib` folder of your Flink distribution

✓ Sampling rate

增加 `job.metric.sample.rate` 配置项，抽样率用于控制 metric 的计算频率  
Add the configuration of `job.metric.sample.rate`, sampling rate is used to control how often certain metrics are computed

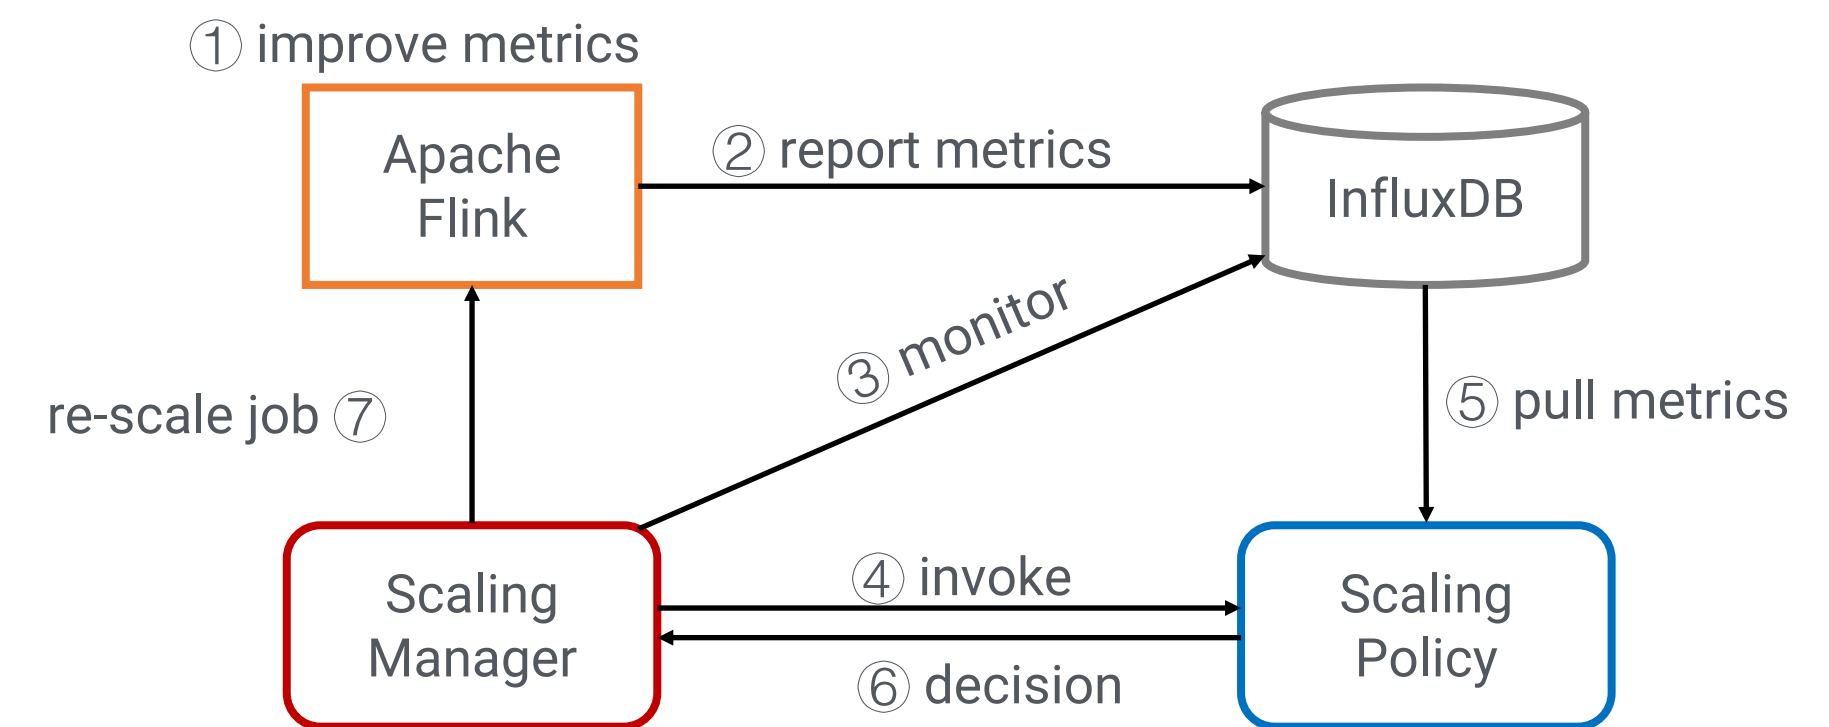


# 生产实践

Production Practices

## ③ monitor

Scaling Manager 监控 InfluxDB  
monitor the InfluxDB with scaling manager

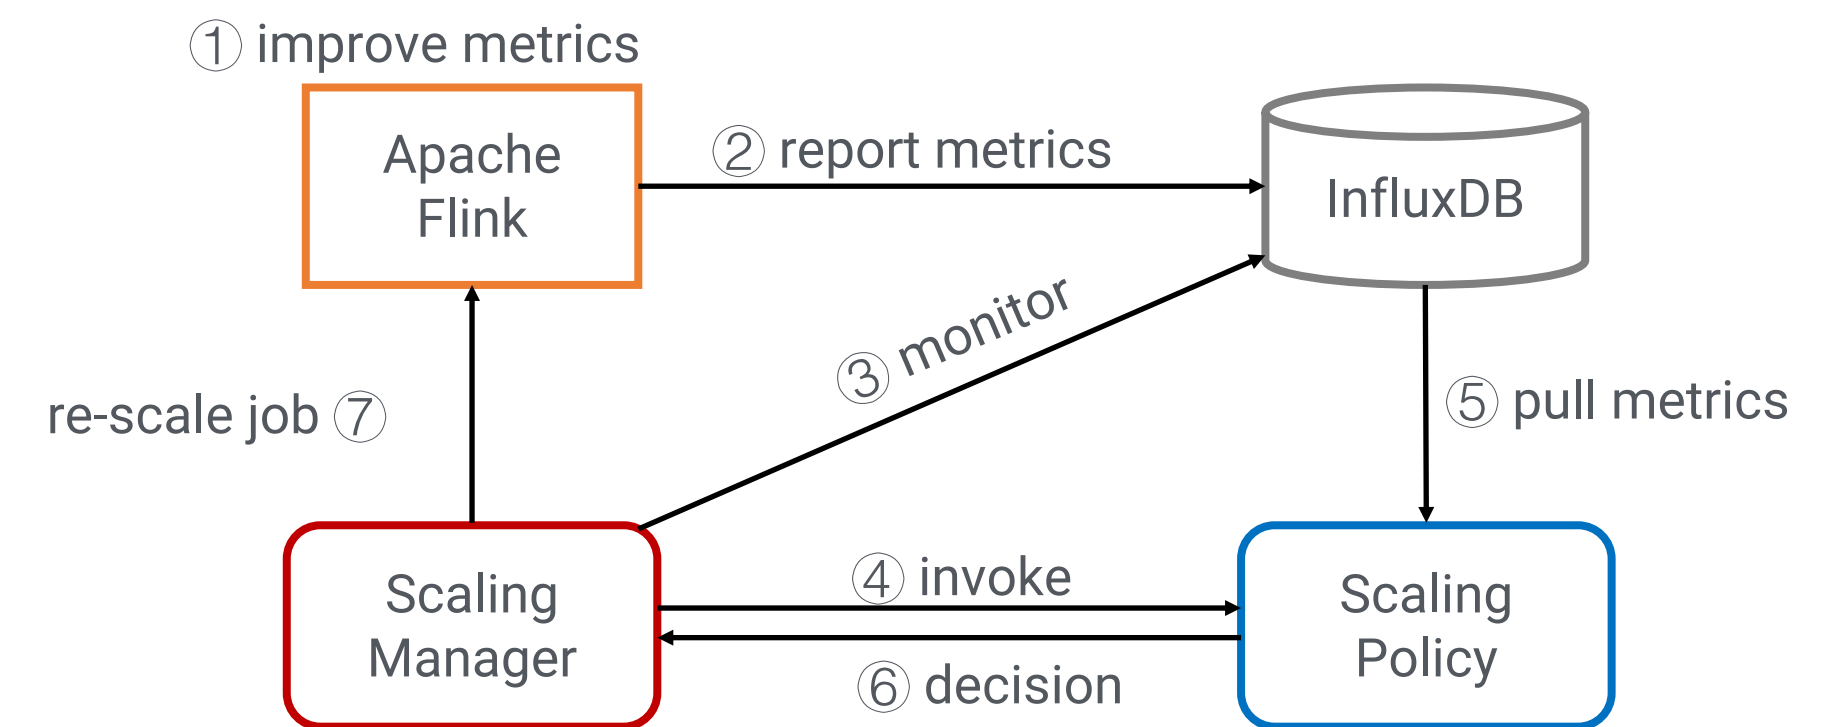


# 生产实践

## Production Practices

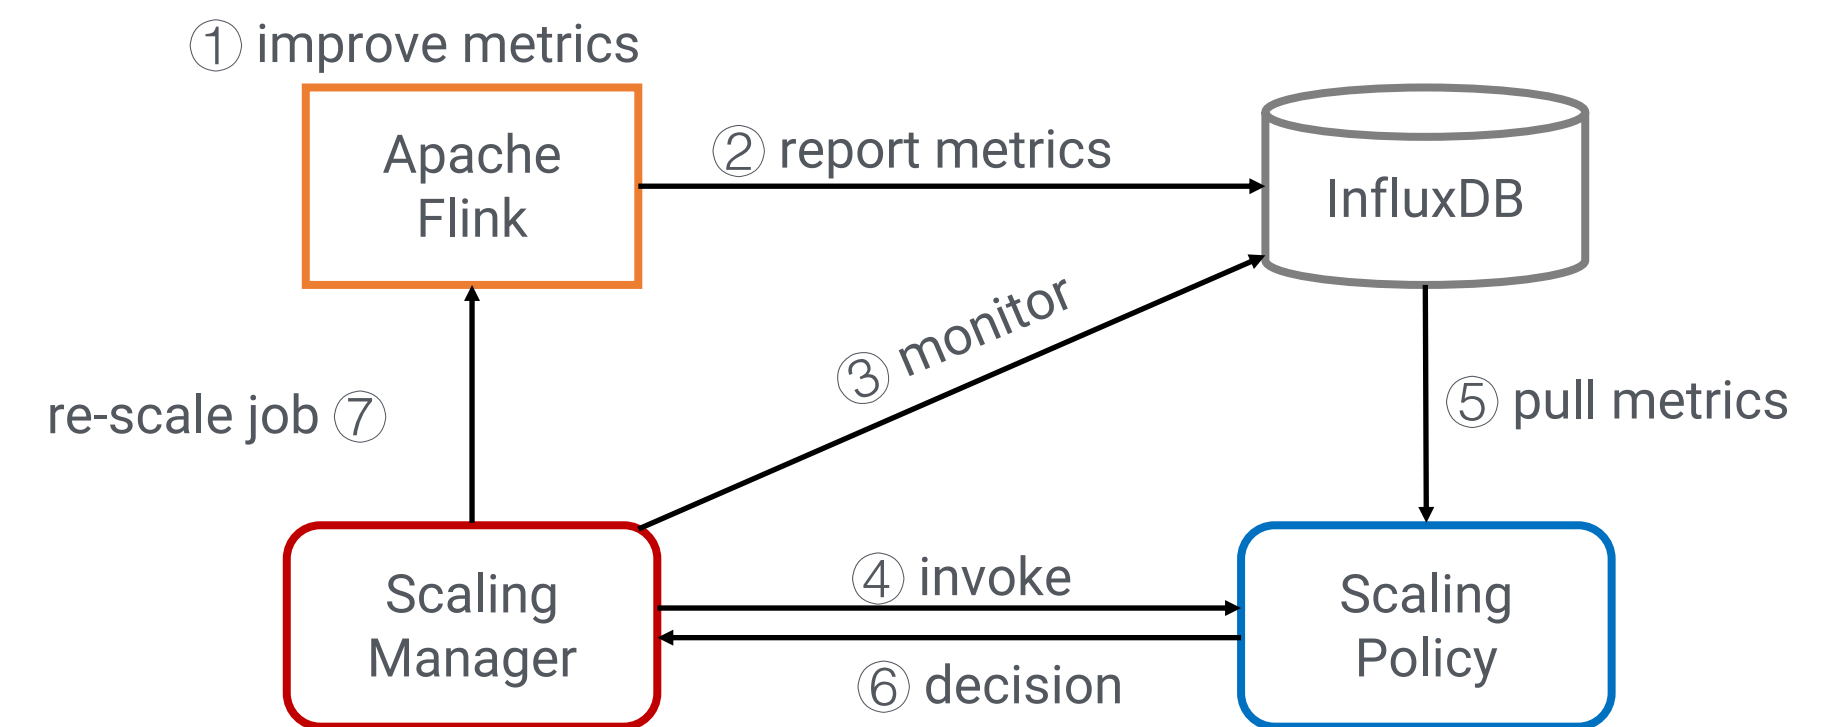
### ④ invoke

当 metrics 可用时，调用策略 Scaling Policy  
invoke the scaling policy when metrics are available



# 生产实践

## Production Practices



### ⑤ pull metrics

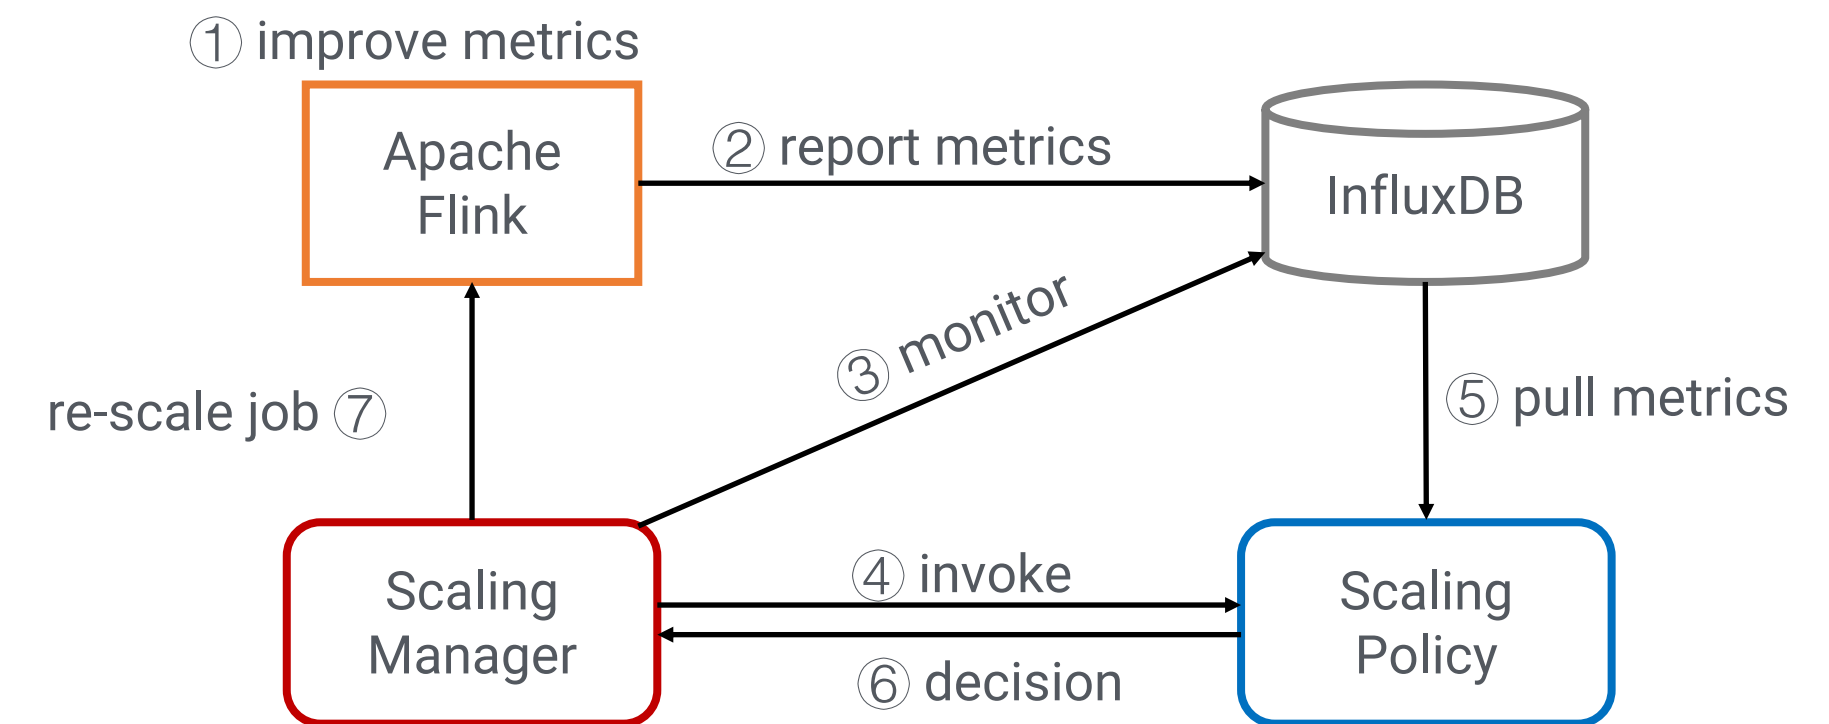
- ✓ taskmanager\_job\_task\_operator\_numRecordsIn
- ✓ taskmanager\_job\_task\_operator\_numRecordsInPerSecond

检查数据倾斜并告警  
check data skew and alarm



# 生产实践

## Production Practices



### ⑤ pull metrics

- ✓ taskmanager\_job\_task\_operator\_KafkaConsumer\_topic\_partition\_0\_committedOffsets
- ✓ taskmanager\_job\_task\_operator\_KafkaConsumer\_topic\_partition\_0\_latestOffsets

计算源速率

compute source rate

检查 kafka 堆积

check the kafka lag

自定义源速率

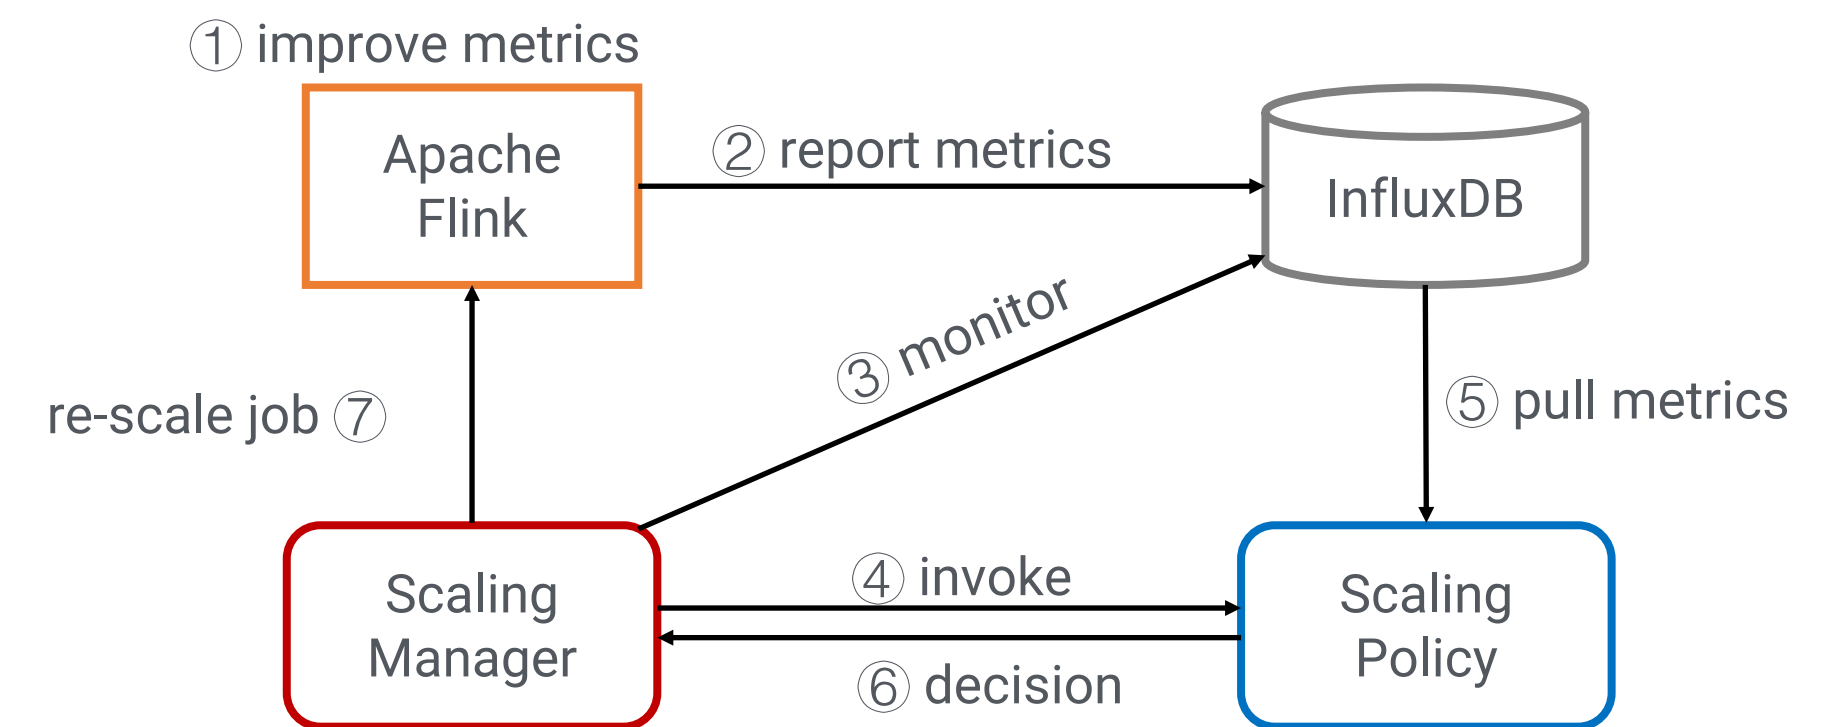
Custom source rate

# 生产实践

## Production Practices

### ⑤ pull metrics

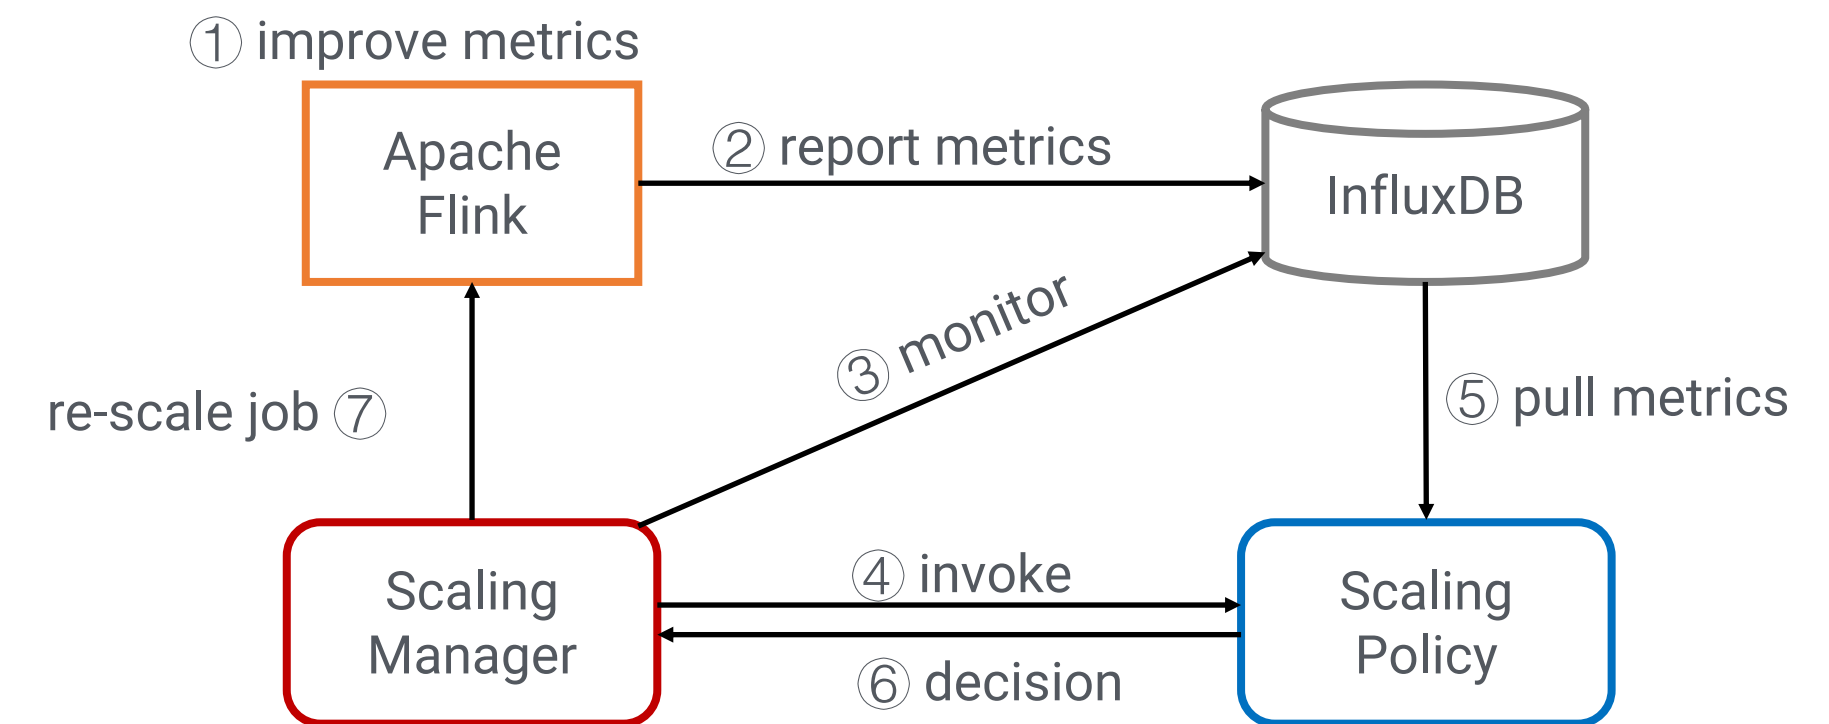
- ✓ taskmanager\_job\_task\_operator\_deserializationDuration
- ✓ taskmanager\_job\_task\_operator\_processingDuration
- ✓ taskmanager\_job\_task\_operator\_serializationDuration
- ✓ taskmanager\_job\_task\_operator\_waitingDuration



计算每个算子速率  
compute each operator's rate

# 生产实践

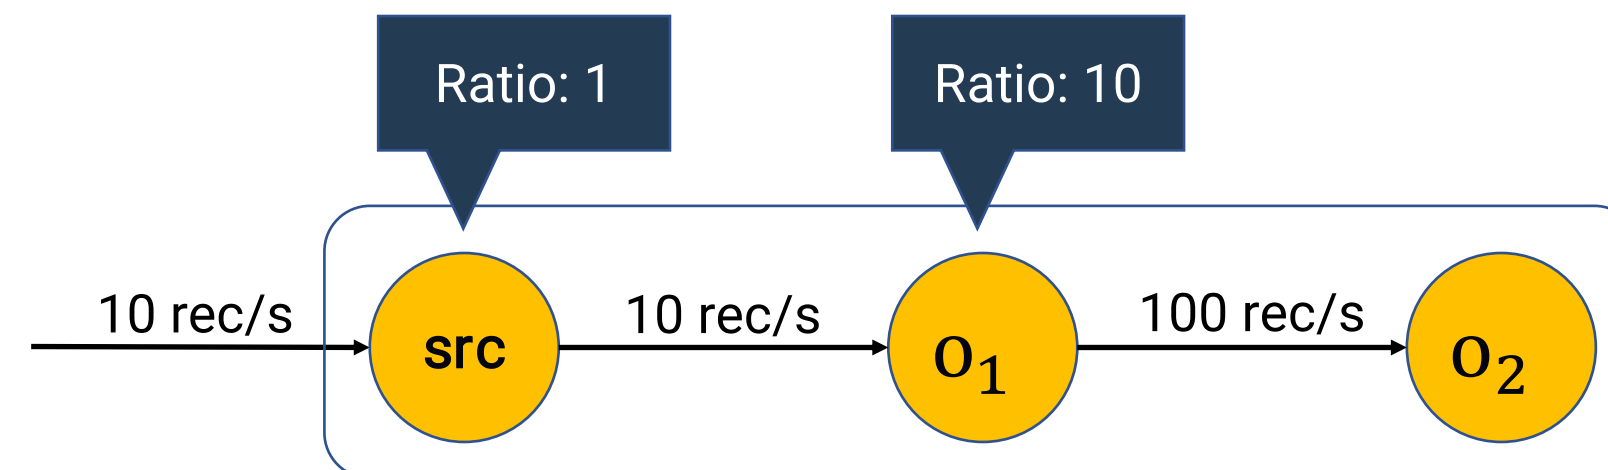
## Production Practices



### ⑤ pull metrics

- ✓ `taskmanager_job_task_operator_numRecordsIn`
- ✓ `taskmanager_job_task_operator_numRecordsOut`

计算生产消费比  
compute the pc ratio



$$\text{Optimal parallelism for } o_i = \frac{\text{aggregated account of upstream ops} * \text{the pc ratio}}{\text{average true processing rate of } o_i} \times (1 + \text{scaling factor})$$



# 生产实践

## Production Practices

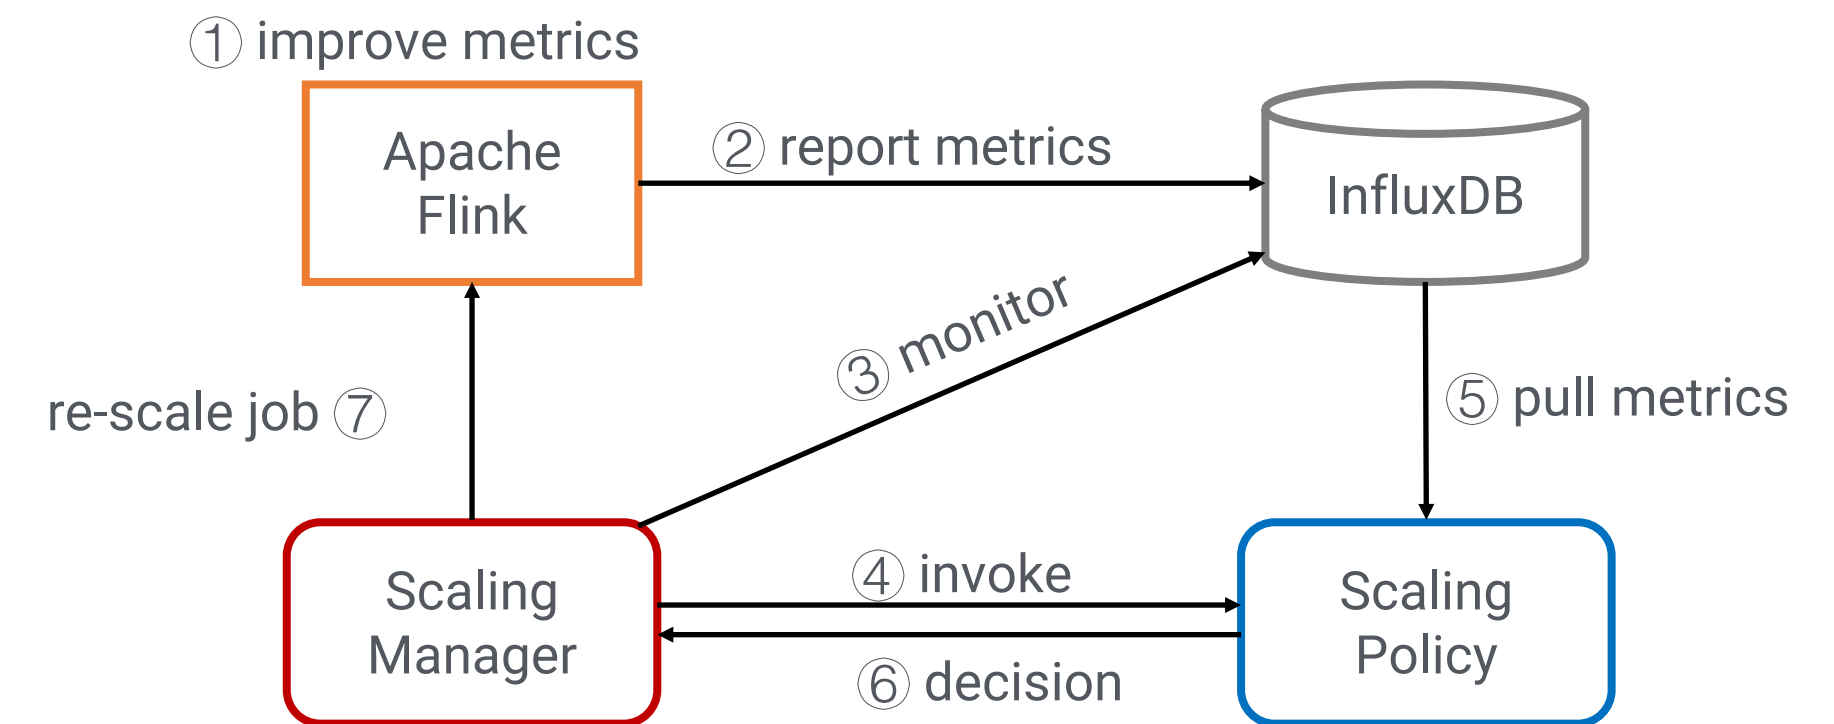
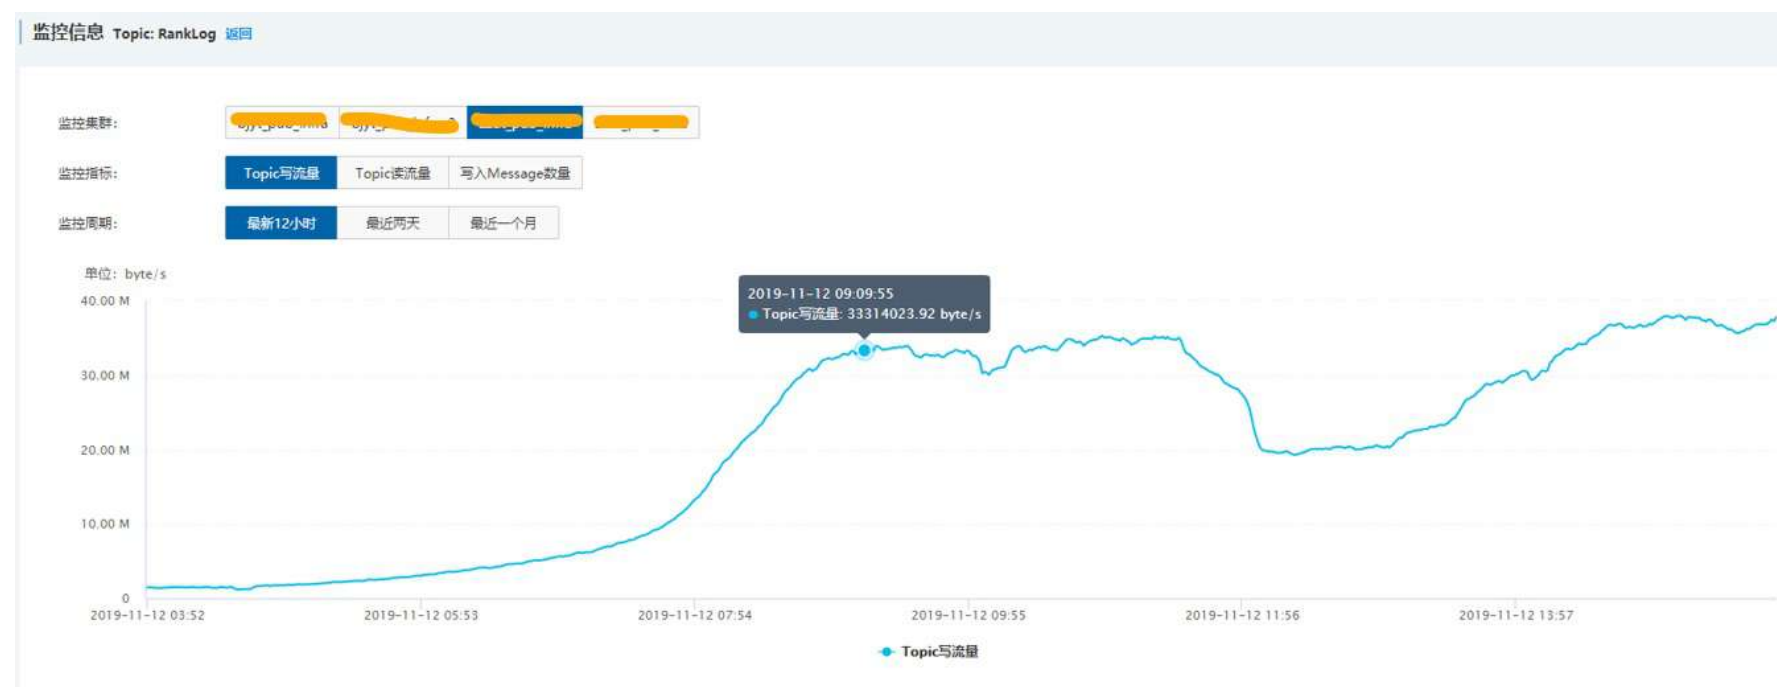
### ⑥ decision

忽略微小改动  
ignore minor changes

源速率趋势图

+

调整幅度

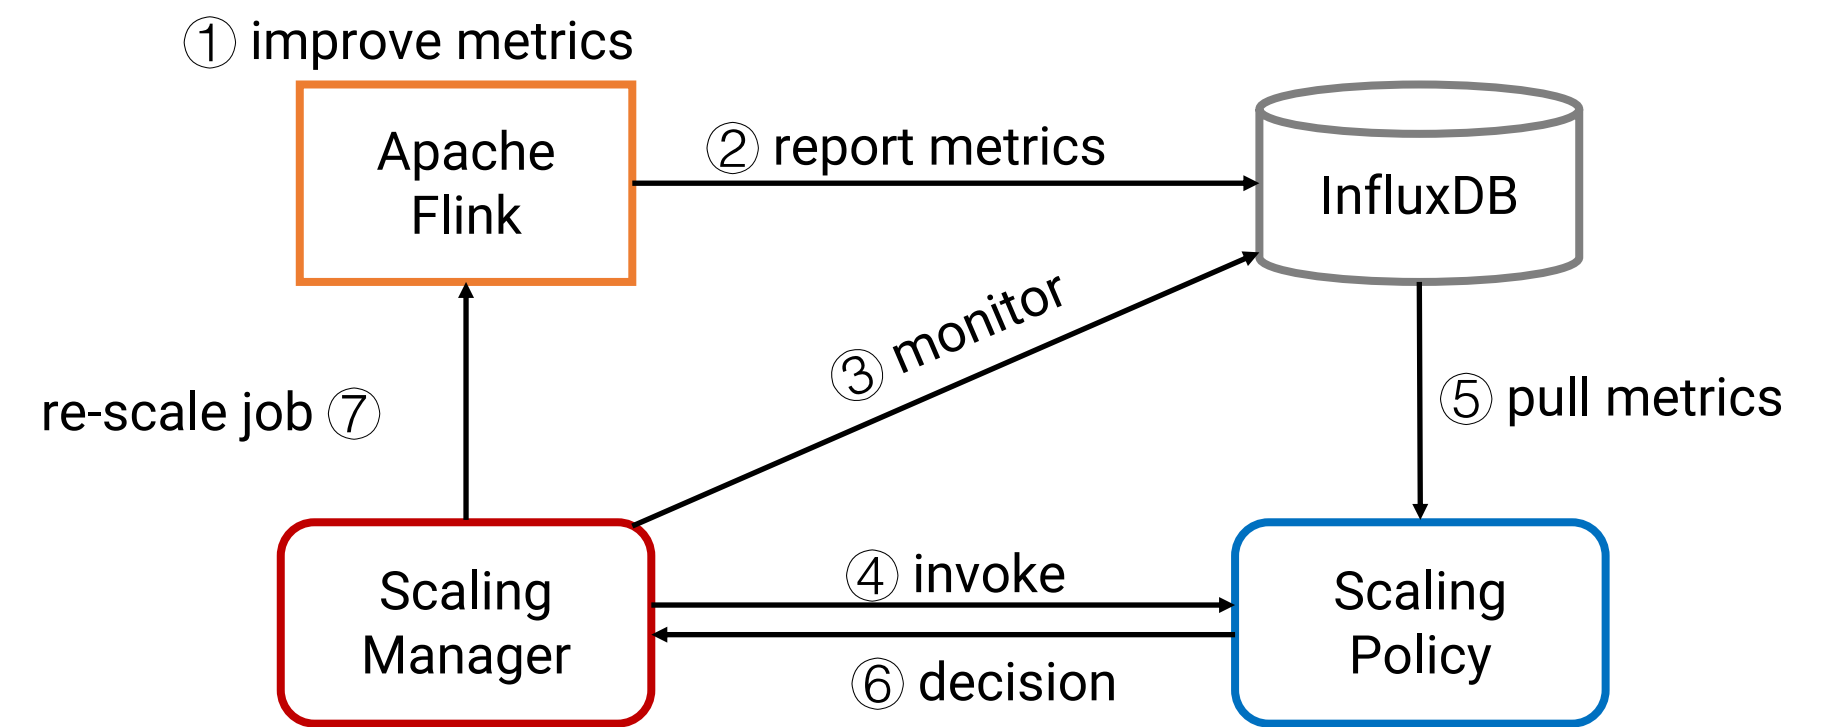


if(before < after && after \* 1.0 / before > (1 + configuration))  
return true;

if(before > after && before \* 1.0 / after > (1 + configuration))  
return true;

# 生产实践

## Production Practices



### ⑦ re-scale job

触发 Savepoint，暂停作业，改造 JobGraph，并使用新的并行度重新部署。  
Flink triggers a Savepoint, halts the computation, reconstructs jobgraph and redeploys it with the new parallelism.

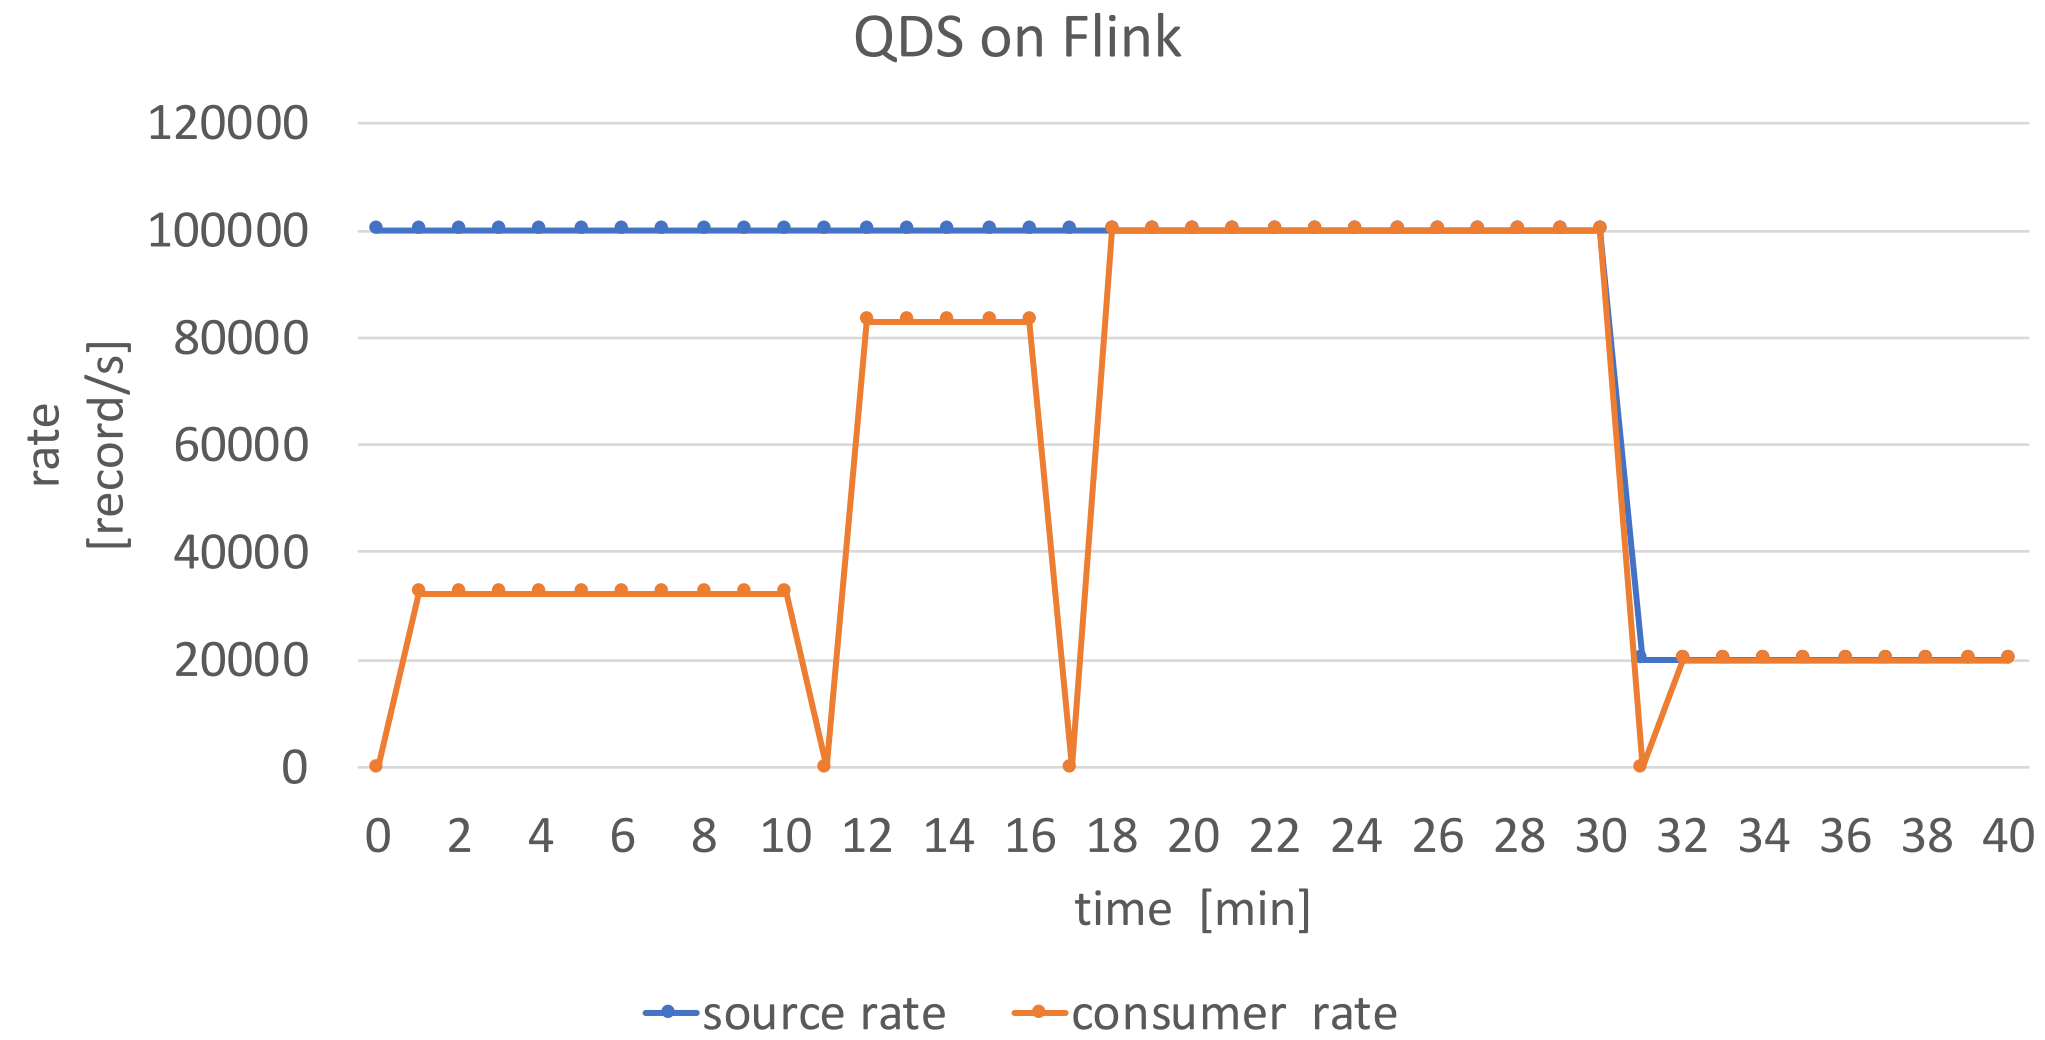
不释放 container 资源，减少作业停止时间。  
Do not release container resources, reduce the job stop time.

✓ yarn-session 模式：调整 resourcemanager.taskmanager-timeout 配置项  
yarn-session mode: adjust the configuration of resourcemanager.taskmanager-timeout

✓ per-job 模式：增加 rescaling 命令  
per-job mode: add the command of rescaling

# 总结

## Summary

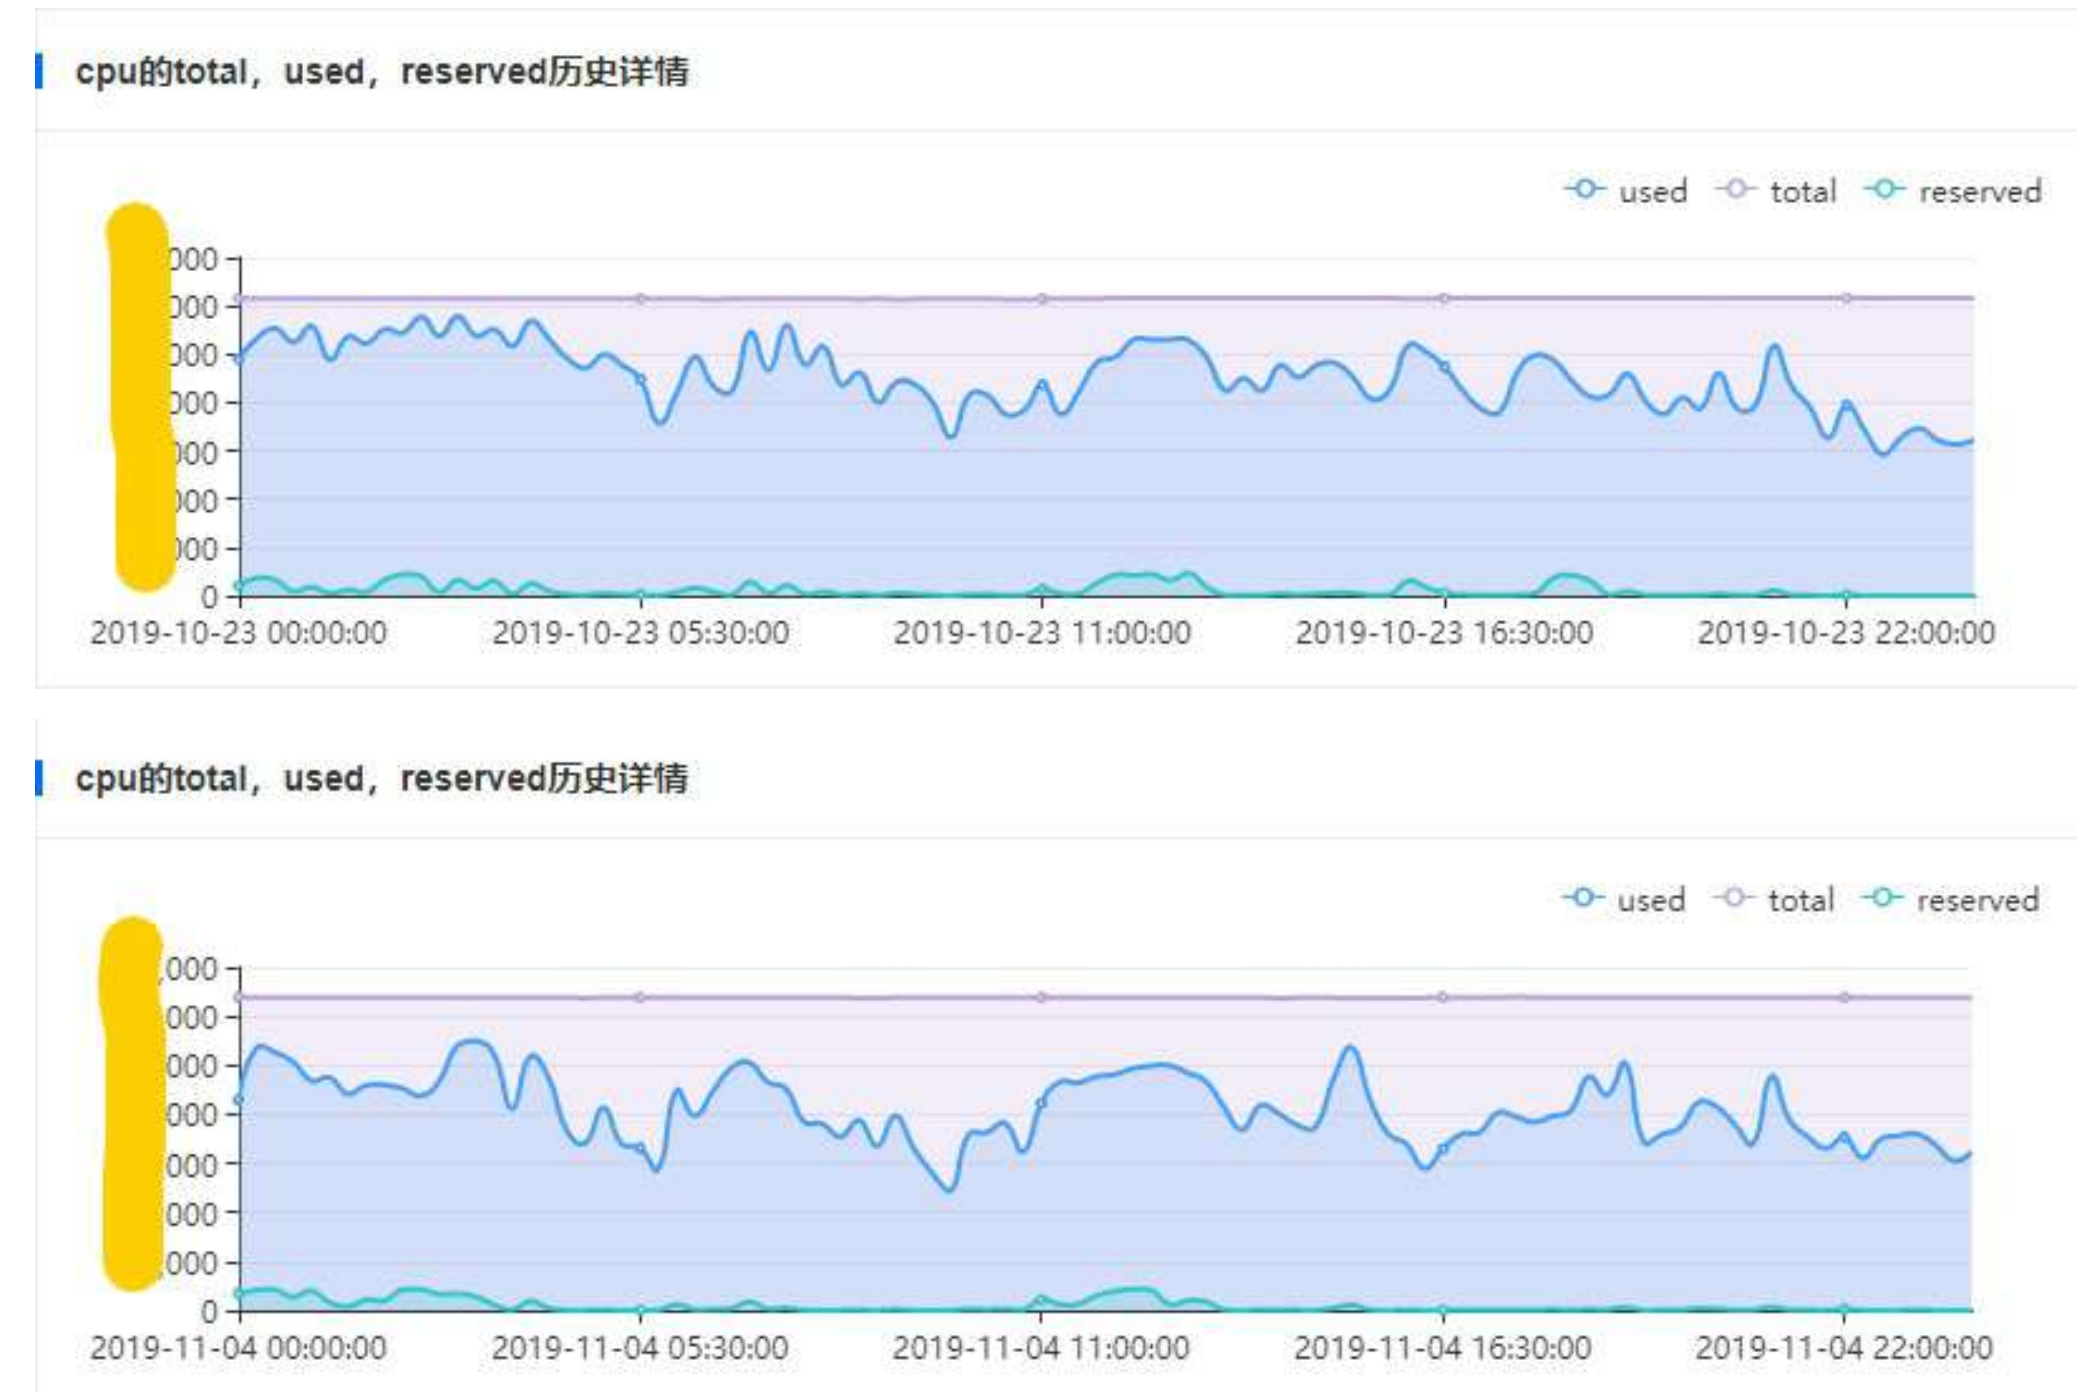


X 同学：可以更多的去关注业务逻辑本身了。

Student X: we can focus on the business logic itself.

实践表明：任务最多三次调整即可达到收敛，集群的资源利用率也提高了。

Practice shows that stream processing took up to three steps to converge, and the resource utilization of the yarn cluster is improved.





# 展望

Future Work

---

04

# 展望

## Future Work

- 进一步提高准确性，“关键路径”分析模型  
Further improve accuracy, critical path analysis model
- 进一步提高稳定性，更快的重新配置机制  
Further improve stability, a faster, more dynamic reconfiguration mechanism
- 回馈社区  
Contribute back to the community

# 致谢

Acknowledgement

Kalavri V, Liagouris J, Hoffmann M, Dimitrova D, Forshaw M, Roscoe T.  
Three steps is all you need: fast, accurate, automatic scaling decisions for  
distributed streaming dataflows.  
OSDI '18.



**THANKS**